

CRITICBENCH: Evaluating Large Language Models as Critic

Tian Lan^{1*} Wenwei Zhang^{2*} Chen Xu³ Heyan Huang¹
 Dahua Lin² Kai Chen^{2†} Xian-Ling Mao^{1†}

¹School of Computer Science and Technology, Beijing Institute of Technology

²Shanghai AI Laboratory

³School of Medical Technology, Beijing Institute of Technology

Abstract

Critique ability are crucial in the scalable oversight and self-improvement of Large Language Models (LLMs). While many recent studies explore the critique ability of LLMs to judge and refine flaws in generations, how to comprehensively and reliably measure the critique abilities of LLMs is under-explored. This paper introduces CRITICBENCH, a novel benchmark designed to comprehensively and reliably evaluate **four key critique ability dimensions of LLMs: feedback, comparison, refinement and meta-feedback**. CRITICBENCH encompasses nine diverse tasks, each assessing the LLMs’ ability to critique responses at varying levels of quality granularity. Our extensive evaluations of open-source and closed-source LLMs reveal intriguing relationships between the critique ability and tasks, response qualities, and model scales. Datasets, resources and evaluation toolkit for CRITICBENCH will be publicly released at <https://github.com/open-compass/CriticBench>.

1 Introduction

Critique ability, *i.e.* the ability to judge and refine flaws, is the key to self-improvement (Yuan et al., 2024; Xu et al., 2023) of Large Language Models (LLMs). Powerful critique ability of LLMs also facilitates a more robust framework, *i.e.*, scalable oversight (Bowman et al., 2022; Saunders et al., 2022), for ensuring the AI systems growing in scale and capability remain aligned with human-desired outcomes and ethical standards.

While numerous studies (Gou et al., 2024; Pan et al., 2023; Yuan et al., 2024; Xu et al., 2023) have explored the critique abilities to enhance the performance of LLMs on some specific tasks, a reliable and comprehensive framework for assessing the critique capabilities is under-explored. Most existing works (Wang et al., 2023c; Li et al., 2024; Cui et al.,

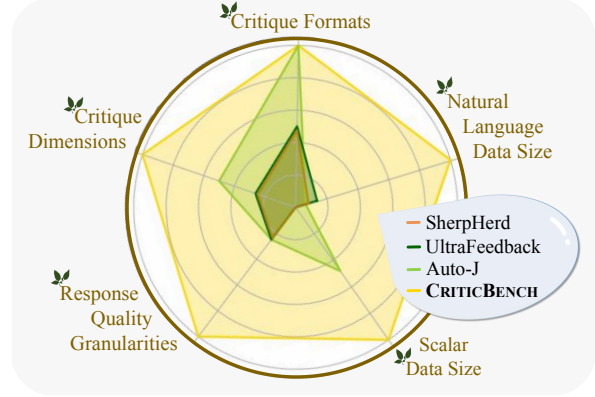


Figure 1: Comparison between CRITICBENCH and existing benchmarks for critique evaluation. More details about the comparison are in Appendix D.

2023) mainly evaluate a singular aspect of critique ability, limited in the critique ability dimensions, application tasks and response qualities. Besides, they frequently utilize GPT-4 (OpenAI, 2023) for direct evaluation of critique quality, not to mention that **GPT-4 is not entirely reliable** (Wang et al., 2023c; Zhang et al., 2024). These issues significantly impede the in-depth analysis and further improvement of critique capabilities.

To address the above problems, we introduce **a diverse and novel benchmark**, CRITICBENCH, aiming to comprehensively and reliably measure the critique capabilities of LLMs. Specifically, to ensure the comprehensiveness of the benchmark, CRITICBENCH evaluates the critique ability of LLMs under nine diverse and common task scenarios, including NLP tasks (summary, translation, and question answering), reasoning (math with program-of-thought (Chen et al., 2023a) and chain-of-thought (Wei et al., 2023)), coding (code with and without execution¹ (Liu et al., 2023a)) and alignment (general chat and harmlessness). We collect model responses of varying quality levels to be critiqued using various open-source and closed-

* Equal contributions

† Corresponding author

¹We hereinafter refer to “code w/. execution” as “Code-Exec” and “code w/o. execution” as “CodeNE”.

source LLMs of different scales and capabilities.

The measurement of critique ability is formulated into four dimensions: **feedback**, **correction**, **comparison**, and **meta-feedback** (*i.e.*, feedback of feedback), completely covering the capability requirements of the self-improvement (Yuan et al., 2024) and scalable oversight (Saunders et al., 2022). Since both the scalar-based and textual formats of critique are commonly used in these scenarios, CRITICBENCH evaluate the critiques in both formats, equipped with objective (Fu et al., 2023; Li et al., 2023) and subjective (Li et al., 2024; Wang et al., 2023c) evaluations, respectively.

To improve the reliability of subjective scoring for textual critiques, we employ **high-quality human-annotated critiques** as references to assist GPT-4, which has been proven effective in enhancing the quality of GPT-4 assessments (Sun et al., 2024a; Kim et al., 2024). We contribute a human-in-the-loop pipeline (Liu et al., 2023b) to efficiently annotate critiques, first generated by GPT-4 and then reviewed and refined by human experts.

Overall, CRITICBENCH exhibits significant advantages over previous benchmarks on critique evaluation (Figure 1), showing great diversity in response quality granularity, critique formats, critique dimensions, and data size, allowing deeper analysis of the LLMs’ critique capabilities. Extensive evaluations of numerous LLMs reveal several intriguing phenomena as below:

- The critique difficulty varies with tasks, where responses of mathematics and coding tasks are challenging for feedback and comparison.
- The critique difficulty correlates with the critique dimensions, where meta-feedback is more challenging than feedback.
- The quality of critiques is inversely related to the response qualities to be critiqued, whereas high-quality responses are more difficult for LLMs to critique sufficiently.
- The critique ability obeys the scaling law, *i.e.*, LLMs with grander scale and general capabilities demonstrate better critique abilities.

There is also a promising trend that some open-source LLMs (Bai et al., 2023; Team, 2023; DeepSeek-AI, 2024) are approaching state-of-the-art closed-source LLMs in critique capabilities.

These phenomena promote an in-depth understanding of LLM’s critique ability, and we hope the discoveries could spur future research in this field.

```
##### Task Input and Response  $R_a, R_b$ 
 $I$ : Please give me a recipe for pizza.
 $R_a$ : Here is a recipe for pizza: ...
 $R_b$ : I don't know any recipe for pizza ...
##### Single-feedback  $F_s(I, R_b)$  for  $R_b$ 
Generated response fails to fulfill user's query ...
Following suggestions could be considered ...
Score for Response  $R_B$ : 1/10
##### Correction  $CR(I, R_b, F_s)$  with  $F_s$ 
Here is the revision for response B:
Of course! To make a pizza, the following ingredients
are needed: ... The steps for cooking are as follows: ...
##### Comparison-based feedback  $F_c(I, R_a, R_b)$ 
Response A is more helpful than B ... On the contrary,
response B only claims it cannot solve this input ...
Preference Response: A
##### Meta-feedback  $F_s(I, R_b, F_s)$ 
Score for Feedback  $F_s$ : 8/10 The feedback point out
the errors in response B ... It also provide suggestions ...
```

Figure 2: Cases of key concepts. **Red** and **brown** represent scalar-based and textual critiques, respectively.

2 Preliminaries

We first formally define the key concepts and their corresponding notions in CRITICBENCH. Figure 2 shows a specific case to understand these concepts.

Task Input (I) and Response (R) represent the user queries and generations produced by LLMs.

Critique aims to analyze and refine the generated responses. Formally, this paper studies the critique capabilities in four dimensions: (1) single-feedback or feedback F_s involves textual analysis and a quality score, also noted as $F_s(I, R)$. Good feedback should not only find flaws but also provide helpful suggestions for refinement (Saunders et al., 2022); (2) correction or refinement CR aims to revise responses with or without feedback. Previous evaluations (Li et al., 2024; Cui et al., 2023) overlook this dimension, although it is an inevitable step when letting the model improve itself (Bai et al., 2022b); (3) comparison or comparison-based feedback F_c contains a textual critique and a preference label for a pair of responses (R_a, R_b). Such an ability is commonly used in the current RLAIIF pipeline that relies on LLMs to label the preferences of responses (Bai et al., 2022b); (4) Meta-feedback $F_s(I, R, F_s)$, *i.e.*, the feedback of feedback itself (Saunders et al., 2022), involves a rating score reflecting the quality of F_s and corresponding textual analysis, which is a higher-level critique dimension. Such an ability is necessary to improve the critique ability (Yuan et al., 2024).

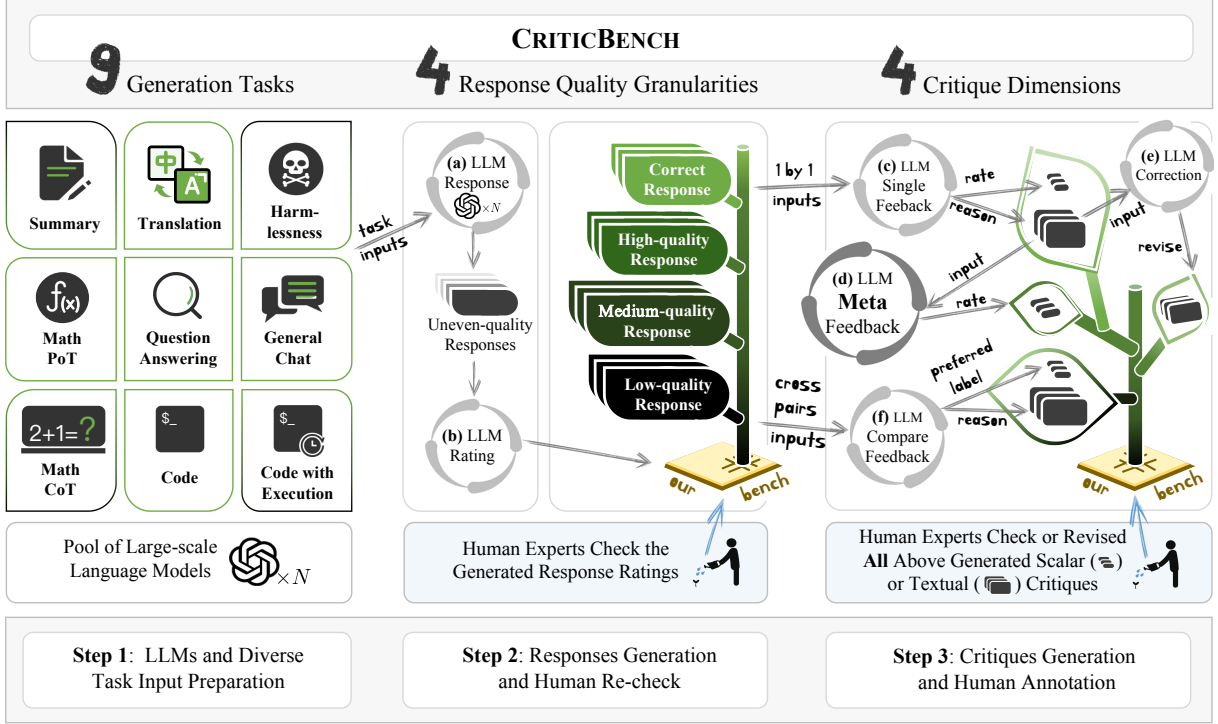


Figure 3: The data construction pipeline for our proposed CRITICBENCH. **Step 1**: 9 tasks and numerous LLMs are prepared. **Step 2**: LLMs are employed to generate responses, which are then meticulously reviewed by human experts. **Step 3**: Critiques are generated by LLMs with strong critique ability, and human experts annotate them.

3 CRITICBENCH Construction

Given the challenge of crafting scalar-based and textual critiques from scratch, we construct CRITICBENCH using a human-in-the-loop data construction pipeline as shown in Figure 3.

3.1 Instruction Collection

Instructions for 9 distinct tasks are collected to evaluate critique capabilities comprehensively (Step 1 in Figure 3). Specifically, the benchmark includes three representative classical language tasks: summary (Stiennon et al., 2020), translation (Specia et al., 2020), and question-answering (Mihaylov et al., 2018). Since a popular application of LLMs is to serve as a chatbot, where alignment is important to ensure the safe application of LLMs, we collect instructions from general chat scenarios (Li et al., 2023) and harmlessness cases (Bai et al., 2022b) to evaluate the LLMs’ critique ability for alignment. Furthermore, the reasoning and code capabilities are also fundamental for augmenting LLMs as agents (Wang et al., 2023a), another important and promising application of LLMs. Thus, we also collect instructions for math reasoning with chain-of-thought and program-of-thought, and coding with and without execution results. To ensure the difficulty of CRITICBENCH, we only collect

coding and math reasoning questions that some 70B LLMs cannot correctly answer. For each task, we collect around 100 instructions from the test sets of some widely-used benchmark datasets to ensure the instruction quality and avoid data contamination. More details are in Appendix A.

3.2 Response Generation

For each collected I in each task, LLMs of different scales and capabilities (e.g., Vicuna (Chiang et al., 2023) and WizardLM (Xu et al., 2024)) are employed to generate responses with flaws, which naturally form responses of various qualities (Step 2 (a) in Figure 3). The complete list of LLMs are in Appendix B. To identify the quality of these responses efficiently, GPT-4 is utilized to initially assign quality ratings ranging from 1 to 7 (Step 2 (b) in Figure 3) then let human annotators meticulously review and adjust these scores, which are used in the objective evaluation in the feedback dimension (Section 4.1). Subsequently, three responses with distinct quality differences for each I are chosen based on their human-varified quality scores, including low-, medium-, and high-quality responses (noted as R_{low} , R_{med} , R_{high} , respectively).

Moreover, golden or correct responses are collected for each task input I , which are proven chal-

lenging for critiques (Wang et al., 2023c; Zhang et al., 2024). We use GPT-4 to generate correct responses using ground-truth rationales or codes as hints for coding and mathematical tasks. Since executions leak information about response quality, correct responses are not collected for the Code-Exec task. In tasks beyond coding and mathematics, GPT-4 is prompted to refine its past generations, given its feedback during multiple turns, and the last revision is collected as golden response.

After sampling task inputs and their corresponding responses R_{low} , R_{med} , R_{high} , we collect the feedback to be evaluated for the meta-feedback dimension by utilizing 4 LLMs that are known powerful for critiques (Step 3 (d) in Figure 3): (1) GPT-4; (2) GPT-3.5-turbo; (3) Auto-J-13B (Li et al., 2024); (4) UltraCM-13B (Cui et al., 2023).

The average annotated scores for collected responses with varying qualities on 9 tasks can be found in Appendix C.

3.3 Reference Critique Generation

After collecting instructions and the corresponding responses, we collect reference critiques on these responses to make the subjective evaluation more reliable, with the assistance of GPT-4.

Feedback and Correction Due to the connection between feedback and correction dimension, GPT-4 is utilized to generate feedback and corrections sequentially (Step 3 (c) and (e) in Fig. 3). We observe that responses in mathematical and coding tasks pose significant challenges for critiques. Thus, ground-truth answers are provided for GPT-4 as references to generate high-quality feedback and corrections. Then, these feedback and corrections are revised by human annotators and used as references in subjective evaluation, respectively.

Comparison-based Feedback To investigate the impact of the quality discrepancy between two responses’ qualities on F_c , given responses $(R_{\text{low}}, R_{\text{med}}, R_{\text{high}})$ of different quality for each instruction I , two pairs are created: $(R_{\text{low}}, R_{\text{high}})$, $(R_{\text{med}}, R_{\text{high}})$. Then we prompt GPT-4 to provide preferences and textual analysis on these pairs (Step 3 (f) in Figure 3), followed by human annotators revising these labels and critiques.

Meta-Feedback Since recent studies (Wang et al., 2023c; Zhang et al., 2024) prove that GPT-4 might incorrectly evaluate the quality of critiques, we ask three human experts to directly provide their Likert scores for generated critiques, ranging from

1 to 7, respectively. Due to the complexity of the meta-feedback dimension, we do not collect its textual critiques and leave it for future research.

Statistics of Human Annotation The average revision rates of textual critiques for the feedback, comparison, and correction are 25.22%, 34.83%, and 48.37%, respectively, indicating that GPT-4’s textual feedback is more likely to be accepted by human annotators, followed by the comparison-based feedback. In contrast, GPT-4’s correction might contain more flaws that need to be corrected. Moreover, the average inner correlation score among annotators is 0.79, indicating solid inner agreement. More details about the statistics of CRITICBENCH and human annotation protocol can be found in Appendix D and E, respectively.

4 Evaluation Metrics

4.1 Objective Evaluation

Feedback and meta-feedback evaluation aim to evaluate the consistency between generated scores and average human annotation scores. It is worthy noting that LLMs are prompted with the annotated reference critiques on the meta-feedback dimension. The widely-used Spearman correlations (Zar, 2005) are computed (Lan et al., 2020; Tao et al., 2017; Liu et al., 2023c; Fu et al., 2023), which ranges from -1 to 1 (normalize to $(-100, 100)$). Higher scores indicate a higher consistency with human judgments.

Comparison-based feedback evaluation assesses the accuracy of LLM in deciding preferences between two responses. It is widely acknowledged that current LLMs exhibit significant **positional bias** (Wang et al., 2023b; Zheng et al., 2023; Zeng et al., 2024), *i.e.*, LLMs tend to prefer responses based on their specific position in the prompt. We implement a rigorous verification process to mitigate the effects of positional bias to evaluate the real capability. Specifically, given responses R_a and R_b to be compared, we obtain the comparison based on two orders, noted as $F_c^a = F_c(R_a, R_b)$ and $F_c^b = F_c(R_b, R_a)$. The objective scores are computed by: $s = \frac{1}{N} \sum_{i=1}^N \mathbb{1}(L(F_c^a, F_c^b))$. $L(F_c^a, F_c^b)^2$ is true if and only if $F_c^a \neq F_c^b$ and F_c^a, F_c^b align with ground-truth preference label. N is the number of test samples.

Correction evaluation is only conducted on mathematics and coding tasks since the revision

² $\mathbb{1} \rightarrow \{0, 1\}$ is the indicator function.

could be easily verified with the ground-truth answers and the test cases. Thus, the objective evaluation metric is implemented as the pass rate: $\frac{N_{\text{Pass}}}{N}$, where N and N_{Pass} are the number of the total samples and passed samples, respectively.

More details about the objective evaluation in CRITICBENCH are in Appendix G.

4.2 Subjective Evaluation

The subjective evaluation aims to evaluate the quality of the generated textual critiques. Since responses in mathematics and coding tasks can be verified, we only conduct subjective evaluations on other 5 tasks for the correction dimension. Following previous works (Wang et al., 2023c; Cui et al., 2023; Li et al., 2024), GPT-4 is leveraged to score the generated critiques. Based on our observations, GPT-4 may not analyze all the flaws effectively. Thus, we prompt GPT-4 with annotated critiques as a reference, facilitating a more reliable subjective evaluation. The case study in Appendix H.1 proves the effectiveness of annotated reference critiques in CRITICBENCH, and more detailed proving is left for future work. It is well-known that LLMs like GPT-4 prefer longer generations during their automatic evaluation, *i.e.*, **length bias** (Wang et al., 2023d; Zeng et al., 2024). However, Figure 5 in Appendix F proves that there is no clue that GPT-4, with our concise and precise reference critiques as input, would give higher scores to longer critiques.

The subjective scores range from 1 to 10. Scores of 1 and 10 denote the worst and best performance, respectively. To effectively measure high-quality critiques, we set the quality score of human-annotated critiques as 8, and the critiques with scores > 8 are better than reference critiques.

4.3 Overall Score

The overall scores of LLMs on these four critique dimensions are calculated to reflect their overall capability. For objective evaluation, the overall scores are the average scores over four critique dimensions, ranging from 0 to 100. Since Spearman scores range from -1 to 1 , they are first normalized into $(0, 100)$ before averaging. For subjective evaluation, since meta-feedback does not require textual critiques, the overall scores are set as the average of feedback, comparison-based feedback, and correction scores, ranging from 1 to 10.

5 Evaluation and Analysis

Extensive experimental results for some representative LLMs on the test set are shown in this section, including 1,010 and 1,214 samples for subjective and objective evaluation, respectively. Due to the page limitation, complete evaluation results of all LLMs on the test and dev set for each task and each critique dimension are placed in Appendix I. Except for the overall analysis of closed-source and open-source LLMs in Section 5.2, we also mainly analyze several intriguing phenomena about the relationship between LLM’s critique ability and some critical factors, including the application tasks (Section 5.3), the difficulty of critique dimensions (Section 5.4), the response quality (Section 5.5), and the scaling law (Section 5.6).

5.1 LLMs to be Evaluated

We extensively evaluate widely used open-source and closed-source LLMs of different sizes on CRITICBENCH to understand the current progress in this field. For closed-source LLMs, we test GPT-4, Claude, Gemini-Pro, PaLM, GPT-3.5-turbo, *etc.* For open-source LLMs, we test different model series so that we could also study scaling laws, including Mistral (Jiang et al., 2024), LLaMA2 (Touvron et al., 2023), Baichuan2 (Baichuan, 2023), Qwen (Bai et al., 2023), InternLM2 (Team, 2023), WizardLM (Xu et al., 2024), Vicuna (Chiang et al., 2023), Yi, and DeepSeek (DeepSeek-AI, 2024), *etc.* Critique-tuned LLMs are also included (Cui et al., 2023; Li et al., 2024) to analyze the effectiveness of the critique corpus. Moreover, we show CRITICBENCH can also be used to evaluate reward models that are fundamental component of RLHF (Ouyang et al., 2022): (1) UltraRM-13B (Cui et al., 2023); (2) Ziya-7B (Wang et al., 2022); (3) SteamSHP (Ethayarajh et al., 2022).

5.2 Overall Analysis of LLMs

As shown in Table 1 and Table 2, GPT-4 significantly outperforms other LLMs on most critique dimensions, approaching the reference critiques (a score of 8). Besides, open-source LLMs are approaching state-of-the-art closed-source LLMs. For example, InternLM2-20B-Chat surpasses GPT-3.5-turbo on overall subjective scores ($6.20 > 5.89$). Surprisingly, InternLM2-7B-Chat achieves comparable performance (5.66) with much bigger LLMs on the subjective evaluation, like Qwen-72B-Chat and DeepSeek-67B-Chat, but losing to them on the objective evaluation. Critique-tuned LLMs

Model	Feedback					Correction						Comp-Feedback			Avg.
	Low	Med.	High	Corr.	Avg.	Low	Med.	High	Avg.	EAvg.	SAvg.	Easy	Hard	Avg.	
Closed-source LLM															
GPT-4-turbo	8.39	8.08	7.86	6.07	7.84	7.91	7.54	7.63	7.69	7.86	8.00	8.30	7.99	8.04	7.86
Claude-instant-1	6.49	5.76	5.29	6.17	5.88	7.74	7.69	7.73	7.72	6.67	6.66	6.16	5.66	5.76	6.45
GPT-3.5-turbo	5.80	4.73	4.63	6.04	5.21	7.61	7.34	7.63	7.55	4.95	5.50	5.15	4.84	4.92	5.89
Open-source LLM (Larger than 30B)															
Qwen-72B-Chat	6.29	5.28	5.01	5.92	5.57	7.56	7.29	7.51	7.45	6.08	6.18	5.20	5.00	5.02	6.01
DeepSeek-67B-Chat	6.21	5.23	5.35	5.39	5.53	7.48	7.26	7.18	7.30	5.09	6.04	5.08	4.54	4.69	5.84
Mixtral-8x7B-v0.1	5.74	5.14	4.88	5.76	5.31	7.35	7.23	7.40	7.33	5.11	5.93	4.82	4.63	4.62	5.75
Llama2-70B-Chat	3.85	4.22	4.39	4.21	4.12	7.03	7.11	7.17	7.11	5.58	5.54	4.07	3.97	3.95	5.06
Critique-tuned LLM (13B)															
Auto-J-13B	4.87	4.38	4.24	3.10	4.21	-	-	-	-	-	-	4.98	4.57	4.63	-
UltraCM-13B	4.07	3.88	3.07	4.84	4.12	-	-	-	-	-	-	-	-	-	-
Open-source LLM (13B-20B)															
InternLM2-20B-Chat	6.73	5.77	5.68	6.05	6.03	7.35	7.75	7.33	7.48	4.93	5.80	5.40	5.06	5.10	6.20
Qwen-14B-Chat	4.85	4.40	4.56	5.84	4.81	7.28	7.09	7.39	7.25	4.58	4.88	4.00	4.11	3.98	5.35
Llama2-13B-Chat	3.53	3.78	3.94	3.65	3.70	6.92	7.24	7.16	7.11	4.69	4.56	3.52	3.26	3.32	4.71
Open-source LLM (6B-7B)															
InternLM2-7B-Chat	5.51	5.11	4.75	5.82	5.20	7.31	7.01	7.19	7.17	4.66	5.50	4.84	4.54	4.62	5.66
Mistral-7B-ins-v0.2	4.90	4.44	4.46	5.36	4.70	7.19	7.23	7.19	7.20	4.82	5.34	4.52	4.24	4.28	5.39
Qwen-7B-Chat	3.59	3.83	4.24	4.96	4.05	6.26	6.34	6.53	6.38	4.10	4.09	3.57	3.48	3.47	4.63
DeepSeek-7B-Chat	3.33	3.22	3.36	4.22	3.44	6.25	5.66	6.26	6.06	4.09	4.23	3.75	3.56	3.60	4.37
Llama2-7B-Chat	3.26	3.65	3.52	3.49	3.44	6.34	6.44	6.02	6.02	4.24	4.00	3.14	3.31	3.21	4.22
Average for Each Quality Granularity															
Average	5.14	4.76	4.66	5.11	4.89	7.17	7.08	7.15	7.12	5.16	5.48	4.78	4.55	4.58	5.56

Table 1: Subjective evaluation results (1-10) on the test set of CRITICBENCH. LLM’s evaluation results are shown on samples with different qualities. Light gray denotes the worst and second-worst performance, while dark gray denotes the best and second-best performance. Best average scores are highlighted in **bold**. Last column (Avg.) contains the overall scores. **EAvg** and **SAvg** denote correction with empty and LLM’s self-feedback, respectively. Since partial models are not optimized for correction and comparison dimensions, their results are not recorded.

outperform Llama-2 series models while still lagging behind some powerful open-source LLMs, like the InternLM2 series. Besides, reward models like UltraRM-13B achieve impressive performance in scoring the quality of responses, significantly outperforming GPT-3.5-turbo and Claude models.

5.3 Relationship with Tasks

Effective critiques usually require domain knowledge and understanding of given tasks. Thus, we also analyze the relationship between critique ability and tasks in Table 3.

Feedback LLMs achieve much higher scores in the first five tasks than on math reasoning and coding tasks (average subjective scores $4.92 > 3.19$ and average objective scores $27.29 > 18.53$), indicating that math and code tasks are more challenging for critiques. Responses in CodeExec are slightly easier to critique than those in CodeNE, implying a positive effect of code execution information for self-improvement on coding tasks.

Comparison-based Feedback Similarly, the first five tasks are much easier to critique than mathematics and coding tasks. Besides, LLMs achieve the best scores on QA, chat, and summary tasks, while average scores on coding tasks are the worst.

Correction Mathematics tasks are more challenging on objective evaluation than coding tasks, and CodeExec are more accessible to revise than those in CodeNE due to richer information in executions. The translation is the most challenging task because professional domain knowledge is required, while harmlessness is the easiest to refine since most LLMs have been trained to avoid harmful generations (Bai et al., 2022b; Lee et al., 2023).

Meta-Feedback LLMs achieve much higher consistency with human judgments on code and mathematics tasks than other 5 tasks (average score $23.06 > 3.12$). LLMs perform much worse on translation and summary tasks. This observation proves that automatically evaluating critiques on mathematics and code tasks is more reliable. Moreover, as shown in Table 2, although GPT-4 achieves a very high correlation with human judgments (62.90) on the meta-feedback dimension, its gap to the average human level (79.03) is 16.13, indicating an apparent divergence between the GPT-4 and human judgments in assessing critique quality, which has also been observed but not quantitatively verified in recent works (Wang et al., 2023c).

Model	Feedback (−100,100)					Correction (%)						Comp-Feedback (%)			Meta-Feedback (−100,100)				Avg.
	Low	Med.	High	Correct	Avg.	Low	Med.	High	Avg.	EAvg.	SAvg.	Easy	Hard	Avg.	Low	Med.	High	Avg.	
Closed-source LLM																			
GPT-4-turbo	53.64	61.82	49.98	16.84	63.54	66.88	69.48	72.75	69.67	60.67	52.17	63.98	53.03	57.33	59.85	66.49	60.59	62.90	72.55
GPT-3.5-turbo	43.90	36.56	28.67	18.96	51.44	66.56	61.04	63.34	64.00	18.67	21.00	51.08	33.33	40.67	41.47	18.84	25.88	28.71	60.83
Claude-instant-1	24.75	24.67	21.78	34.47	42.78	49.90	46.70	55.36	50.00	24.84	20.17	55.91	37.12	44.89	49.55	35.70	32.42	38.89	58.93
Open-source LLM (> 30B)																			
DeepSeek-67B	30.51	21.61	11.56	28.06	42.11	57.71	57.67	50.70	55.00	21.00	21.83	52.69	40.53	45.56	26.95	37.11	34.31	33.60	59.60
Qwen-72B-Chat	25.76	23.74	7.59	36.72	42.64	61.15	47.99	58.09	54.67	14.83	7.83	54.30	36.74	44.00	19.14	37.70	25.11	27.86	58.48
Mixtral-8x7B-v0.1	35.46	39.61	14.09	55.51	51.00	52.81	37.08	44.27	43.34	14.34	14.00	47.85	40.91	43.78	18.84	33.08	28.71	26.66	56.49
Llama2-70B-Chat	21.05	25.79	33.10	20.02	32.79	39.69	38.61	47.22	42.34	8.50	5.33	22.58	20.08	21.11	37.66	28.76	19.81	28.32	48.50
Open-source LLM (13B-20B)																			
InternLM2-20B	46.69	43.53	25.66	19.00	58.61	62.19	37.71	55.48	50.50	5.67	10.33	52.15	39.39	44.67	0.12	-1.43	26.99	3.95	56.61
Llama2-13B-Chat	11.21	17.63	26.22	37.37	30.61	18.23	30.87	22.55	24.67	5.17	4.50	29.03	18.18	22.67	26.64	14.14	50.02	31.02	44.54
Qwen-14B-Chat	-10.28	2.08	16.88	45.21	14.32	33.96	45.21	35.19	38.00	8.33	6.67	16.67	15.15	15.78	15.70	4.96	16.00	10.72	41.58
Critique-tuned LLM (13B)																			
Auto-J-13B	21.16	32.59	32.54	4.11	36.05	-	-	-	-	-	-	53.23	46.59	49.33	-	-	-	-	-
UltraCM-13B	-5.54	7.58	29.97	28.70	21.51	-	-	-	-	-	-	38.17	37.88	38.00	-	-	-	-	-
Open-source LLM (6B-7B)																			
Mistral-7B	33.55	37.16	38.16	25.35	43.66	51.98	38.65	31.62	38.17	11.17	13.00	36.02	21.97	27.88	32.44	31.05	29.84	31.68	50.93
InternLM2-7B	43.23	40.52	13.02	32.95	49.09	49.90	23.37	38.67	36.17	5.67	6.00	33.87	16.67	23.78	8.90	-9.22	20.99	3.17	46.52
DeepSeek-7B	-9.39	-0.51	1.7	18.42	8.26	32.08	21.84	46.48	35.00	5.83	8.17	20.43	18.56	19.33	-7.73	-11.77	21.98	-0.31	39.57
Llama2-7B-Chat	-0.54	12.03	11.53	36.86	20.81	16.98	23.40	21.11	21.00	5.50	2.83	6.45	4.55	5.33	5.92	-0.07	15.75	5.67	34.89
Qwen-7B-Chat	-11.7	-28.88	-23.87	9.97	-8.09	30.52	28.96	34.97	32.33	6.67	1.33	6.99	4.17	5.33	5.48	11.43	24.14	11.73	34.87
Reward Models																			
UltraRM-13B	47.42	29.33	39.81	18.06	52.33	-	-	-	-	-	-	65.05	47.35	54.67	-	-	-	-	-
Ziya-7B	15.84	11.98	13.42	17.10	25.81	-	-	-	-	-	-	48.39	34.09	40.00	-	-	-	-	-
Average for Each Quality Granularity																			
Average	21.93	23.10	20.62	26.51	35.75	46.04	40.58	45.19	43.66	14.44	13.01	39.73	29.80	33.90	22.73	19.78	28.84	22.97	50.99

Table 2: Objective evaluation on the test set. Last column (Avg.) contains the overall scores.

Tasks	$F_s(I, R)$		$F_c(I, R_a, R_b)$		CR		$F_s(F_s)$
	Sub.	Obj.	Sub.	Obj.	Sub.	Obj.	Obj.
Translation	4.43	31.14	3.78	18.28	5.31	-	-2.93
Chat	5.09	20.60	4.97	32.60	5.66	-	1.80
QA	5.20	30.75	5.05	27.67	6.42	-	13.50
Summary	4.76	28.93	4.63	37.12	5.99	-	0.54
Harmless.	5.12	25.04	3.97	19.35	7.51	-	2.71
Avg.	4.92	27.29	4.48	27.00	6.18	-	3.12
MathCoT	3.55	22.56	2.80	12.42	-	29.36	19.63
MathPoT	3.35	27.80	3.05	14.98	-	24.98	22.73
CodeExec	3.07	13.38	2.74	7.72	-	32.20	25.50
CodeNE	2.77	10.37	2.80	10.33	-	29.50	24.38
Avg.	3.19	18.53	2.85	11.36	-	29.01	23.06

Table 3: Average performance of all LLMs on 9 tasks. Two **Avg.** rows represent the average scores of the first 5 tasks and the last 4 tasks, respectively.

5.4 Difficulty of Critique Dimensions

As shown in last rows of Table 1 and Table 2, the comparison-based feedback presents the highest level of difficulty (4.58), followed by the feedback dimension (average score 4.89), proving that comparison mostly requires accurate analysis of both responses, which is more complex than the single feedback dimension. As a higher-level critique dimension, meta-feedback is more challenging than the feedback dimension ($22.97 < 35.75$).

Moreover, the correction is the easiest critique dimension with our annotated feedback (average score 7.12). To understand how feedback affects the quality of refinements, we replace the annotated feedback with Empty feedback (**EAvg.** column) and LLM’s Self-feedback (**SAvg.** column)

for the correction dimension, and several conclusions can be summarized: (1) most LLMs achieve the best performance with our annotated feedback in CRITICBENCH, proving the effectiveness of our human annotations; (2) high-quality feedback is beneficial to refinement on subjective evaluation ($\text{Avg}(7.12) > \text{SAvg}(5.48) > \text{EAvg}(5.16)$), while low-quality feedback might negatively affect refinement on objective evaluation ($\text{SAvg}(13.01) < \text{EAvg}(14.44) < \text{Avg}(43.66)$). This observation proves that LLMs struggle in self-improvement on challenging mathematics and coding tasks, aligning with recent findings (Huang et al., 2024); (3) since GPT-4’s average score on subjective evaluation of the feedback dimension is slightly better than annotated feedback $8.09 > 8$ (first row in Table 13), it is reasonable that GPT-4 achieves best performance with its feedback. On the contrary, GPT-4’s average feedback score on mathematics and coding tasks is $7.53 < 8$ (first row in Table 13), leading to worse performance on the objective evaluation ($52.17 < 69.67$).

5.5 Relationship with Response Quality

Effective self-improvement needs to identify which kinds of responses are hard to critique. Thus, we analyze the relationship between response quality and critique ability given the Table 1 and Table 2. Most LLMs perform worst on high-quality responses for the feedback dimension, indicating that

subtle flaws in these responses pose a significant challenge. In contrast, correct and medium-quality responses are the easiest to critique on the subjective and objective evaluation, respectively. Besides, GPT-4’s critique ability on correct responses 6.07 is much worse than other responses. This observation also quantitatively confirms recent findings that GPT-4 might modify better answers worse (Stechly et al., 2023; Zhang et al., 2024; Huang et al., 2024). For the correction critique dimension, most LLMs perform worst on medium-quality responses while achieving the best performance on high-quality and low-quality responses. For the comp-feedback critique dimension, it is reasonable that samples with distinct qualities are much easier for critics than the harder samples, *i.e.*, two responses with close quality. For the meta-feedback dimension, most LLMs achieve higher consistency with human judgments on high- and low-quality responses than medium-quality responses³, indicating the difficulty of measuring critiques on medium-quality responses.

5.6 Relationship with Scaling Law

Recent works (Wei et al., 2022) have shown that as the scale of models gradually increases, their capability across various tasks also improves. It is also a crucial question whether critique ability obeys this scaling law. As shown in Tables 1 and 2, most LLMs’ critique ability (objective and subjective evaluation) also improves steadily as the number of parameters increases, as observed in the Qwen series, Llama-2 series, DeepSeek series, Mistral series, and InternLM2 series LLMs. For example, as the parameters of Qwen series LLMs increase from 7B to 72B, the overall subjective scores exhibit a progressive increase from 4.63 to 6.01, and overall objective scores increase from 34.87 to 58.48.

6 Related Work

Self-improvement by Critique Ability Critique ability has been widely used for LLM’s self-improvement in the following two stages: (1) **Inference stage**: through prompting LLMs to generate feedback and refinements sequentially, LLMs could iteratively improve their generations (Pan et al., 2023; Madaan et al., 2023; Zhang et al., 2024; Saunders et al., 2022; Yao et al., 2023; Fernandes et al., 2023; Huang et al., 2024; Saunders et al., 2022); (2) **Training stage**: the refinements could

be used for the supervised fine-tuning (SFT) or preference learning (RLHF) to improve LLM. They are also widely used in the RLAIIF pipeline (Bowman et al., 2022; Saunders et al., 2022; Yuan et al., 2024; Bai et al., 2022b; Xu et al., 2023). Meanwhile, minor human principles are still beneficial to LLM self-alignment. (Sun et al., 2023, 2024b).

Self-improvement highly relies on LLM’s critique ability. However, there are few studies for measuring critique ability comprehensively, especially the correction and meta-feedback dimensions, impeding the in-depth understanding of the bottleneck of self-improvement. This paper proposes the CRITICBENCH to fill this gap.

Improvement of Critique Ability Open-source LLMs are proven to have poor critique ability (Ye et al., 2023; Wang et al., 2023c). Recently, numerous works have been proposed to improve Llama-2 models by fine-tuning on critiques generated by GPT-4 (Li et al., 2024; Cui et al., 2023; Kim et al., 2024; Ke et al., 2023). However, their evaluation is limited in the task scenarios, critique dimensions, critique formats, and test data size.

Application of Critique Ability So far, critique ability is widely used for automatic evaluation of response quality, frequently utilizing GPT-4 (OpenAI, 2023; Liu et al., 2023c; Fu et al., 2023) or critique-tuned LLMs (Ke et al., 2023; Li et al., 2024), whereas they are proved unstable for some complex task scenarios (Wang et al., 2023c; Zhang et al., 2024). Besides, evaluating open-source LLMs’ critique ability is overlooked in previous works (Li et al., 2024; Cui et al., 2023). Our work not only proposes a systematic benchmark but also comprehensively evaluates closed-source and open-source LLMs on critique ability (Touvron et al., 2023; DeepSeek-AI, 2024; Xu et al., 2024; Chiang et al., 2023; Team, 2023; Bai et al., 2023), showing several intriguing phenomena.

7 Conclusion and Future Work

In this paper, we introduce a systematic and comprehensive benchmark, named CRITICBENCH, specifically developed to assess the critique capabilities of LLMs. Furthermore, we conducted an extensive evaluation of both open-source and closed-source LLMs on CRITICBENCH. The in-depth analysis shed light on the critique capabilities of these models concerning critique dimensions, tasks, response qualities, and model sizes, significantly promoting the development of LLM’s

³Responses here denote the generations R to be critiqued, not the generated critiques F_g in the meta-feedback dimension.

self-improvement and scalable oversight. Such an analysis also pinpoints potential pitfalls in self-improvement and scalable oversight of LLMs, especially when the LLM becomes more powerful.

In the future, we plan to enhance our benchmark and evaluation protocol in several key areas: (1) Broadening the scope to include more tasks, like tool-using; (2) Extending the benchmark to encompass other languages, like Chinese; (3) Improving the subjective evaluation protocol to allow for more fine-grained analysis, particularly verifying the claims made in generated feedback.

8 Limitations

CRITICBENCH mainly utilizes the advanced GPT-4 model for subjective evaluation. Despite integrating high-quality annotated reference critiques to guide GPT-4 toward more accurate assessments, it’s essential to acknowledge that the model’s evaluations may not always align perfectly with human judgment. For example, although GPT-4 achieves strong consistency with average human judgments of 62.90 (Table 2), the gap between GPT-4 and the average human level is 16.13, showing an apparent divergence. While LLMs have yet to reach the level of precision of human annotation - they currently represent the most effective approach for balancing the trade-offs between evaluation cost and quality. It is still a significant challenge to accurately and automatically evaluate critiques across all scenarios. Recognizing this, we aim to address these issues in our future work by progressively refining our benchmark and evaluation protocols.

9 Ethical Considerations

Most task inputs in CRITICBENCH are collected from publicly available datasets, free from any possible harm toward individuals or groups. Moreover, humans carefully select and process the responses and critiques generated by LLMs to secure privacy and confidentiality. No personal identification information is involved. However, it should be noted that the task input, responses, and critiques in the Anthropic-HHH dataset (Bai et al., 2022a) of the harmlessness task contain harmful materials and hate speech. Despite the risks involved, it is essential to disclose this research fully, and materials in the Anthropic-HHH dataset have been widely used for safety research in the LLM community.

References

- Aida Amini, Saadia Gabriel, Shanchuan Lin, Rik Koncel-Kedziorski, Yejin Choi, and Hannaneh Hajishirzi. 2019. [MathQA: Towards interpretable math word problem solving with operation-based formalisms](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 2357–2367, Minneapolis, Minnesota. Association for Computational Linguistics.
- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. 2021. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*.
- Jinze Bai, Shuai Bai, Yunfei Chu, Zeyu Cui, Kai Dang, Xiaodong Deng, Yang Fan, Wenbin Ge, Yu Han, Fei Huang, Binyuan Hui, Luo Ji, Mei Li, Junyang Lin, Runji Lin, Dayiheng Liu, Gao Liu, Chengqiang Lu, Keming Lu, Jianxin Ma, Rui Men, Xingzhang Ren, Xuancheng Ren, Chuanqi Tan, Sinan Tan, Jianhong Tu, Peng Wang, Shijie Wang, Wei Wang, Shengguang Wu, Benfeng Xu, Jin Xu, An Yang, Hao Yang, Jian Yang, Shusheng Yang, Yang Yao, Bowen Yu, Hongyi Yuan, Zheng Yuan, Jianwei Zhang, Xingxuan Zhang, Yichang Zhang, Zhenru Zhang, Chang Zhou, Jingren Zhou, Xiaohuan Zhou, and Tianhang Zhu. 2023. Qwen technical report. *arXiv preprint arXiv:2309.16609*.
- Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, Nicholas Joseph, Saurav Kadavath, Jackson Kernion, Tom Conerly, Sheer El-Showk, Nelson Elhage, Zac Hatfield-Dodds, Danny Hernandez, Tristan Hume, Scott Johnston, Shauna Kravec, Liane Lovitt, Neel Nanda, Catherine Olsson, Dario Amodei, Tom Brown, Jack Clark, Sam McCandlish, Chris Olah, Ben Mann, and Jared Kaplan. 2022a. [Training a helpful and harmless assistant with reinforcement learning from human feedback](#).
- Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, Carol Chen, Catherine Olsson, Christopher Olah, Danny Hernandez, Dawn Drain, Deep Ganguli, Dustin Li, Eli Tran-Johnson, Ethan Perez, Jamie Kerr, Jared Mueller, Jeffrey Ladish, Joshua Landau, Kamal Ndousse, Kamile Lukosuite, Liane Lovitt, Michael Sellitto, Nelson Elhage, Nicholas Schiefer, Noemi Mercado, Nova DasSarma, Robert Lasenby, Robin Larson, Sam Ringer, Scott Johnston, Shauna Kravec, Sheer El Showk, Stanislav Fort, Tamera Lanham, Timothy Telleen-Lawton, Tom Conerly, Tom Henighan, Tristan Hume, Samuel R. Bowman, Zac Hatfield-Dodds, Ben Mann, Dario Amodei, Nicholas Joseph, Sam McCandlish, Tom Brown, and Jared Kaplan. 2022b. [Constitutional ai: Harmlessness from ai feedback](#).

- Baichuan. 2023. [Baichuan 2: Open large-scale language models](#). *arXiv preprint arXiv:2309.10305*.
- Yonatan Bisk, Rowan Zellers, Ronan Le Bras, Jianfeng Gao, and Yejin Choi. 2020. Piqa: Reasoning about physical commonsense in natural language. In *Thirty-Fourth AAAI Conference on Artificial Intelligence*.
- Samuel R. Bowman, Jeeyoon Hyun, Ethan Perez, Edwin Chen, Craig Pettit, Scott Heiner, Kamilė Lukošiušė, Amanda Askell, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, Christopher Olah, Daniela Amodei, Dario Amodei, Dawn Drain, Dustin Li, Eli Tran-Johnson, Jackson Kernion, Jamie Kerr, Jared Mueller, Jeffrey Ladish, Joshua Landau, Kamal Ndousse, Liane Lovitt, Nelson Elhage, Nicholas Schiefer, Nicholas Joseph, Noemí Mercado, Nova DasSarma, Robin Larson, Sam McCandlish, Sandipan Kundu, Scott Johnston, Shauna Kravec, Sheer El Showk, Stanislav Fort, Timothy Telleen-Lawton, Tom Brown, Tom Henighan, Tristan Hume, Yuntao Bai, Zac Hatfield-Dodds, Ben Mann, and Jared Kaplan. 2022. [Measuring progress on scalable oversight for large language models](#).
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan, Scott Gray, Nick Ryder, Mikhail Pavlov, Alethea Power, Lukasz Kaiser, Mohammad Bavarian, Clemens Winter, Philippe Tillet, Felipe Petroski Such, Dave Cummings, Matthias Plappert, Fotios Chantzis, Elizabeth Barnes, Ariel Herbert-Voss, William Hebgen Guss, Alex Nichol, Alex Paino, Nikolas Tezak, Jie Tang, Igor Babuschkin, Suchir Balaji, Shantanu Jain, William Saunders, Christopher Hesse, Andrew N. Carr, Jan Leike, Josh Achiam, Vedant Misra, Evan Morikawa, Alec Radford, Matthew Knight, Miles Brundage, Mira Murati, Katie Mayer, Peter Welinder, Bob McGrew, Dario Amodei, Sam McCandlish, Ilya Sutskever, and Wojciech Zaremba. 2021. [Evaluating large language models trained on code](#).
- Wenhu Chen, Xueguang Ma, Xinyi Wang, and William W. Cohen. 2023a. [Program of thoughts prompting: Disentangling computation from reasoning for numerical reasoning tasks](#).
- Wenhu Chen, Ming Yin, Max Ku, Pan Lu, Yixin Wan, Xueguang Ma, Jianyu Xu, Xinyi Wang, and Tony Xia. 2023b. [TheoremQA: A theorem-driven question answering dataset](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 7889–7901, Singapore. Association for Computational Linguistics.
- Wei-Lin Chiang, Zhuohan Li, Zi Lin, Ying Sheng, Zhanghao Wu, Hao Zhang, Lianmin Zheng, Siyuan Zhuang, Yonghao Zhuang, Joseph E. Gonzalez, Ion Stoica, and Eric P. Xing. 2023. [Vicuna: An open-source chatbot impressing gpt-4 with 90%* chatgpt quality](#).
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. 2021. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*.
- Ganqu Cui, Lifan Yuan, Ning Ding, Guanming Yao, Wei Zhu, Yuan Ni, Guotong Xie, Zhiyuan Liu, and Maosong Sun. 2023. [Ultrafeedback: Boosting language models with high-quality feedback](#).
- DeepSeek-AI. 2024. [Deepseek llm: Scaling open-source language models with longtermism](#). *arXiv preprint arXiv:2401.02954*.
- Kawin Ethayarajh, Yejin Choi, and Swabha Swayamdipta. 2022. [Understanding dataset difficulty with V-usable information](#). In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 5988–6008. PMLR.
- Patrick Fernandes, Daniel Deutsch, Mara Finkelstein, Parker Riley, André Martins, Graham Neubig, Ankush Garg, Jonathan Clark, Markus Freitag, and Orhan Firat. 2023. [The devil is in the errors: Leveraging large language models for fine-grained machine translation evaluation](#). In *Proceedings of the Eighth Conference on Machine Translation*, pages 1066–1083, Singapore. Association for Computational Linguistics.
- Jinlan Fu, See-Kiong Ng, Zhengbao Jiang, and Pengfei Liu. 2023. [Gptscore: Evaluate as you desire](#).
- Zhibin Gou, Zhihong Shao, Yeyun Gong, yelong shen, Yujia Yang, Nan Duan, and Weizhu Chen. 2024. [CRITIC: Large language models can self-correct with tool-interactive critiquing](#). In *The Twelfth International Conference on Learning Representations*.
- Jie Huang, Xinyun Chen, Swaroop Mishra, Huaixiu Steven Zheng, Adams Wei Yu, Xinying Song, and Denny Zhou. 2024. [Large language models cannot self-correct reasoning yet](#). In *The Twelfth International Conference on Learning Representations*.
- Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Ye Jin Bang, Andrea Madotto, and Pascale Fung. 2023. [Survey of hallucination in natural language generation](#). *ACM Computing Surveys*, 55(12):1–38.
- Albert Q. Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, Gianna Lengyel, Guillaume Bour, Guillaume Lample, L  lio Renard Lavaud, Lucile Saulnier, Marie-Anne Lachaux, Pierre Stock, Sandeep Subramanian, Sophia Yang, Szymon Antoniak, Teven Le Scao, Th  ophile Gervet, Thibaut Lavril, Thomas Wang, Timoth  e Lacroix, and William El Sayed. 2024. [Mixtral of experts](#).

- Pei Ke, Bosi Wen, Zhuoer Feng, Xiao Liu, Xuanyu Lei, Jiale Cheng, Shengyuan Wang, Aohan Zeng, Yuxiao Dong, Hongning Wang, Jie Tang, and Minlie Huang. 2023. [Critiquellm: Scaling llm-as-critic for effective and explainable evaluation of large language model generation](#).
- Seungone Kim, Jamin Shin, Yejin Cho, Joel Jang, Shayne Longpre, Hwaran Lee, Sangdo Yun, Seongjin Shin, Sungdong Kim, James Thorne, and Minjoon Seo. 2024. [Prometheus: Inducing evaluation capability in language models](#). In *The Twelfth International Conference on Learning Representations*.
- Tian Lan, Xian-Ling Mao, Wei Wei, Xiaoyan Gao, and Heyan Huang. 2020. [Pone: A novel automatic evaluation metric for open-domain generative dialogue systems](#). *ACM Trans. Inf. Syst.*, 39(1).
- Harrison Lee, Samrat Phatale, Hassan Mansoor, Thomas Mesnard, Johan Ferret, Kellie Lu, Colton Bishop, Ethan Hall, Victor Carbune, Abhinav Rastogi, and Sushant Prakash. 2023. [Rlaif: Scaling reinforcement learning from human feedback with ai feedback](#).
- Junlong Li, Shichao Sun, Weizhe Yuan, Run-Ze Fan, haizhao, and Pengfei Liu. 2024. [Generative judge for evaluating alignment](#). In *The Twelfth International Conference on Learning Representations*.
- Xuechen Li, Tianyi Zhang, Yann Dubois, Rohan Taori, Ishaan Gulrajani, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. 2023. AlpacaEval: An automatic evaluator of instruction-following models. https://github.com/tatsu-lab/alpaca_eval.
- Wang Ling, Dani Yogatama, Chris Dyer, and Phil Blunsom. 2017. Program induction by rationale generation: Learning to solve and explain algebraic word problems. *ACL*.
- Jiate Liu, Yiqin Zhu, Kaiwen Xiao, QIANG FU, Xiao Han, Yang Wei, and Deheng Ye. 2023a. [RLTF: Reinforcement learning from unit test feedback](#). *Transactions on Machine Learning Research*.
- Xiao Liu, Xuanyu Lei, Shengyuan Wang, Yue Huang, Zhuoer Feng, Bosi Wen, Jiale Cheng, Pei Ke, Yifan Xu, Weng Lam Tam, Xiaohan Zhang, Lichao Sun, Hongning Wang, Jing Zhang, Minlie Huang, Yuxiao Dong, and Jie Tang. 2023b. [Alignbench: Benchmarking chinese alignment of large language models](#).
- Yang Liu, Dan Iter, Yichong Xu, Shuhang Wang, Ruochen Xu, and Chenguang Zhu. 2023c. [G-eval: NLG evaluation using gpt-4 with better human alignment](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 2511–2522, Singapore. Association for Computational Linguistics.
- Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, Shashank Gupta, Bodhisattwa Prasad Majumder, Katherine Hermann, Sean Welleck, Amir Yazdanbakhsh, and Peter Clark. 2023. [Self-refine: Iterative refinement with self-feedback](#).
- Todor Mihaylov, Peter Clark, Tushar Khot, and Ashish Sabharwal. 2018. Can a suit of armor conduct electricity? a new dataset for open book question answering. In *EMNLP*.
- Swaroop Mishra, Arindam Mitra, Neeraj Varshney, Bhavdeep Sachdeva, Peter Clark, Chitta Baral, and Ashwin Kalyan. 2022. Numglue: A suite of fundamental yet challenging mathematical reasoning tasks. *ACL*.
- Rajani Nazneen, Lambert Nathan, Sheon Han, Wang Jean, Nitski Osvald, Beeching Edward, and Tunstall Lewis. 2023. Can foundation models label data like humans? *Hugging Face Blog*. <https://huggingface.co/blog/llm-v-human-data>.
- OpenAI. 2023. [Gpt-4 technical report](#).
- Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. 2022. [Training language models to follow instructions with human feedback](#).
- Liangming Pan, Michael Saxon, Wenda Xu, Deepak Nathani, Xinyi Wang, and William Yang Wang. 2023. [Automatically correcting large language models: Surveying the landscape of diverse self-correction strategies](#).
- William Saunders, Catherine Yeh, Jeff Wu, Steven Bills, Long Ouyang, Jonathan Ward, and Jan Leike. 2022. [Self-critiquing models for assisting human evaluators](#).
- Lucia Specia, Frédéric Blain, Marina Fomicheva, Erick Fonseca, Vishrav Chaudhary, Francisco Guzmán, and André F. T. Martins. 2020. [Findings of the WMT 2020 shared task on quality estimation](#). In *Proceedings of the Fifth Conference on Machine Translation*, pages 743–764, Online. Association for Computational Linguistics.
- Kaya Stechly, Matthew Marquez, and Subbarao Kambhampati. 2023. [GPT-4 doesn’t know it’s wrong: An analysis of iterative prompting for reasoning problems](#). In *NeurIPS 2023 Foundation Models for Decision Making Workshop*.
- Nisan Stiennon, Long Ouyang, Jeff Wu, Daniel M. Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford, Dario Amodei, and Paul Christiano. 2020. Learning to summarize from human feedback. In *NeurIPS*.
- Shichao Sun, Junlong Li, Weizhe Yuan, Ruifeng Yuan, Wenjie Li, and Pengfei Liu. 2024a. [The critique of critique](#).

- Zhiqing Sun, Yikang Shen, Hongxin Zhang, Qinhong Zhou, Zhenfang Chen, David Daniel Cox, Yiming Yang, and Chuang Gan. 2024b. [SALMON: Self-alignment with principle-following reward models](#). In *The Twelfth International Conference on Learning Representations*.
- Zhiqing Sun, Yikang Shen, Qinhong Zhou, Hongxin Zhang, Zhenfang Chen, David Daniel Cox, Yiming Yang, and Chuang Gan. 2023. [Principle-driven self-alignment of language models from scratch with minimal human supervision](#). In *Thirty-seventh Conference on Neural Information Processing Systems*.
- Alon Talmor, Jonathan Herzig, Nicholas Lourie, and Jonathan Berant. 2019. [CommonsenseQA: A question answering challenge targeting commonsense knowledge](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4149–4158, Minneapolis, Minnesota. Association for Computational Linguistics.
- Chongyang Tao, Lili Mou, Dongyan Zhao, and Rui Yan. 2017. [Ruber: An unsupervised method for automatic evaluation of open-domain dialog systems](#). In *AAAI Conference on Artificial Intelligence*.
- InternLM Team. 2023. Internlm: A multilingual language model with progressively enhanced capabilities. <https://github.com/InternLM/InternLM>.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajwal Bhargava, Shruti Bhosale, Dan Bikel, Lukas Blecher, Cristian Canton Ferrer, Moya Chen, Guillem Cucurull, David Esiobu, Jude Fernandes, Jeremy Fu, Wenyin Fu, Brian Fuller, Cynthia Gao, Vedanuj Goswami, Naman Goyal, Anthony Hartshorn, Saghar Hosseini, Rui Hou, Hakan Inan, Marcin Kardas, Viktor Kerkez, Madian Khabsa, Isabel Kloumann, Artem Korenev, Punit Singh Koura, Marie-Anne Lachaux, Thibaut Lavril, Jenya Lee, Diana Liskovich, Yinghai Lu, Yuning Mao, Xavier Martinet, Todor Mihaylov, Pushkar Mishra, Igor Molybog, Yixin Nie, Andrew Poulton, Jeremy Reizenstein, Rashi Rungta, Kalyan Saladi, Alan Schelten, Ruan Silva, Eric Michael Smith, Ranjan Subramanian, Xiaoqing Ellen Tan, Binh Tang, Ross Taylor, Adina Williams, Jian Xiang Kuan, Puxin Xu, Zheng Yan, Iliyan Zarov, Yuchen Zhang, Angela Fan, Melanie Kambadur, Sharan Narang, Aurelien Rodriguez, Robert Stojnic, Sergey Edunov, and Thomas Scialom. 2023. [Llama 2: Open foundation and fine-tuned chat models](#).
- Junjie Wang, Yuxiang Zhang, Lin Zhang, Ping Yang, Xinyu Gao, Ziwei Wu, Xiaoqun Dong, Junqing He, Jianheng Zhuo, Qi Yang, Yongfeng Huang, Xiayu Li, Yanghan Wu, Junyu Lu, Xinyu Zhu, Weifeng Chen, Ting Han, Kunhao Pan, Rui Wang, Hao Wang, Xiaojun Wu, Zhongshen Zeng, Chongpei Chen, Ruyi Gan, and Jiaying Zhang. 2022. Fengshenbang 1.0: Being the foundation of chinese cognitive intelligence. *CoRR*, abs/2209.02970.
- Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan Chen, Jiakai Tang, Xu Chen, Yankai Lin, Wayne Xin Zhao, Zhewei Wei, and Ji-Rong Wen. 2023a. [A survey on large language model based autonomous agents](#).
- Peiyi Wang, Lei Li, Liang Chen, Zefan Cai, Dawei Zhu, Binghuai Lin, Yunbo Cao, Qi Liu, Tianyu Liu, and Zhifang Sui. 2023b. [Large language models are not fair evaluators](#).
- Tianlu Wang, Ping Yu, Xiaoqing Ellen Tan, Sean O’Brien, Ramakanth Pasunuru, Jane Dwivedi-Yu, Olga Golovneva, Luke Zettlemoyer, Maryam Fazel-Zarandi, and Asli Celikyilmaz. 2023c. [Shepherd: A critic for language model generation](#).
- Yizhong Wang, Hamish Ivison, Pradeep Dasigi, Jack Hessel, Tushar Khot, Khyathi Chandu, David Wadden, Kelsey MacMillan, Noah A. Smith, Iz Beltagy, and Hannaneh Hajishirzi. 2023d. [How far can camels go? exploring the state of instruction tuning on open resources](#). In *Thirty-seventh Conference on Neural Information Processing Systems Datasets and Benchmarks Track*.
- Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, Ed H. Chi, Tatsunori Hashimoto, Oriol Vinyals, Percy Liang, Jeff Dean, and William Fedus. 2022. [Emergent abilities of large language models](#). *Transactions on Machine Learning Research*. Survey Certification.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. 2023. [Chain-of-thought prompting elicits its reasoning in large language models](#).
- Can Xu, Qingfeng Sun, Kai Zheng, Xiubo Geng, Pu Zhao, Jiazhan Feng, Chongyang Tao, Qingwei Lin, and Daxin Jiang. 2024. [WizardLM: Empowering large pre-trained language models to follow complex instructions](#). In *The Twelfth International Conference on Learning Representations*.
- Weiwen Xu, Deng Cai, Zhisong Zhang, Wai Lam, and Shuming Shi. 2023. [Reasons to reject? aligning language models with judgments](#).
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*.
- Seonghyeon Ye, Yongrae Jo, Doyoung Kim, Sungdong Kim, Hyeonbin Hwang, and Minjoon Seo. 2023. [Selfee: Iterative self-revising llm empowered by self-feedback generation](#). Blog post.
- Weizhe Yuan, Richard Yuanzhe Pang, Kyunghyun Cho, Sainbayar Sukhbaatar, Jing Xu, and Jason Weston. 2024. [Self-rewarding language models](#).

Jerrold Zar. 2005. *Spearman Rank Correlation*, volume 5.

Zhiyuan Zeng, Jiatong Yu, Tianyu Gao, Yu Meng, Tanya Goyal, and Danqi Chen. 2024. *Evaluating large language models at evaluating instruction following*. In *The Twelfth International Conference on Learning Representations*.

Wenqi Zhang, Yongliang Shen, Linjuan Wu, Qiuying Peng, Jun Wang, Yueting Zhuang, and Weiming Lu. 2024. *Self-contrast: Better reflection through inconsistent solving perspectives*.

Yue Zhang, Yafu Li, Leyang Cui, Deng Cai, Lemao Liu, Tingchen Fu, Xinting Huang, Enbo Zhao, Yu Zhang, Yulong Chen, Longyue Wang, Anh Tuan Luu, Wei Bi, Freda Shi, and Shuming Shi. 2023. *Siren’s song in the ai ocean: A survey on hallucination in large language models*.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. *Judging LLM-as-a-judge with MT-bench and chatbot arena*. In *Thirty-seventh Conference on Neural Information Processing Systems Datasets and Benchmarks Track*.

A Source Data for Different Tasks

The details of selected datasets for 9 tasks are listed in Table 4, covering the well-known NLP tasks (translation, summary, and question answering), reasoning tasks (mathematics and coding), and alignment (general chat and harmlessness). These datasets’ test sets are used for CRITICBENCH construction, avoiding data contamination.

Tasks	Source From Test Data	Num.
Translation	WMT20 MLQE (Specia et al., 2020)	100
Chat	ChatArena Alpaca-Eval (Li et al., 2023)	50 each
QA	OBQA (Mihaylov et al., 2018) CommonQA (Talmor et al., 2019) PIQA (Bisk et al., 2020)	35 each
Harmlessness	HHH (Bai et al., 2022a)	100
Summary	Summ. HF (Stiennon et al., 2020)	100
Math PoT Math CoT	Aqua-RAT (Ling et al., 2017) MathQA (Amini et al., 2019) GSM8K (Cobbe et al., 2021) NumGLUE (Mishra et al., 2022) TheoremQA (Chen et al., 2023b)	20 each
Code w/. exec Code w/o. exec	MBPP (Austin et al., 2021) HumanEval (Chen et al., 2021)	50 each

Table 4: Source of 9 tasks in CRITICBENCH. Most tasks contain diverse samples from multiple test sets.

B List of Used LLMs

Our study uses several LLMs with different capabilities to generate diverse feedback, listed in Table

5. Besides, we also use some critique-tuned LLMs to generate textual feedback, like Auto-J-13B and UltraCM-13B models.

C Likert Score for Responses

Figure 4 demonstrates the discernible performance disparities in responses for each task. Since automatic execution leaks quality information, we do not collect the correct responses for the CodeExec task. Such variation is instrumental in analyzing the impact of response quality on the feedback.

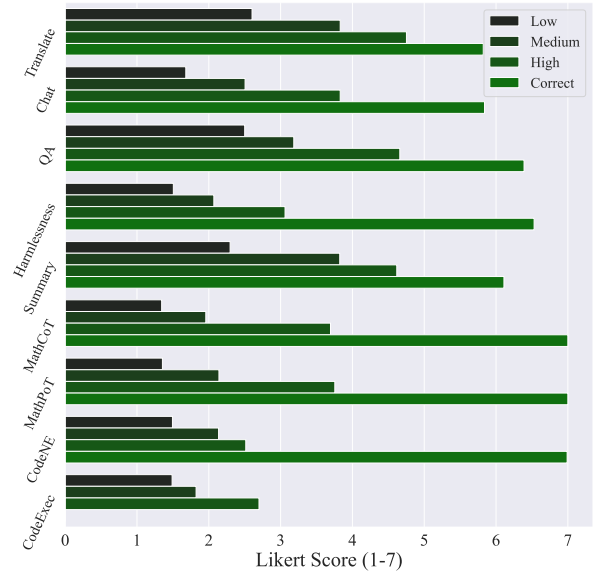


Figure 4: Human annotated Likert scores (1-7) reflecting response quality across tasks. Since execution leaks the information on response quality, correct responses are not collected for the CodeExec task.

D Statics of CRITICBENCH

The comparison between CRITICBENCH and existing benchmarks can be found in Table 6, which proves the advantages of our proposed CRITICBENCH over previous benchmarks for critique evaluation. The number of samples in CRITICBENCH are shown in Table 7. Compared with existing benchmarks for critique evaluation, our proposed CRITICBENCH contains 5,972 samples, leading to a more stable and reliable assessment. Specifically, CRITICBENCH includes 3,608 natural language critique samples (test 1,010 and dev 2598) and 2,892 scalar-based critique samples (test 1,214 and dev 1,688), which human experts well annotate. The scale of CRITICBENCH is much bigger than existing benchmarks for measuring LLM’s critique ability, such as Auto-J (Li et al., 2024), UltraFeedback (Cui et al., 2023), and Shepherd (Wang et al., 2023c).

LLMs	Source
InternLM-7B-8K	https://huggingface.co/internlm/internlm-7b
Qwen-7B-Chat	https://huggingface.co/Qwen/Qwen-7B-Chat
Qwen-14B-Chat	https://huggingface.co/Qwen/Qwen-14B-Chat
Baichuan2-13B	https://huggingface.co/baichuan-inc/Baichuan2-13B-Chat
InternLM-20B	https://huggingface.co/internlm/internlm-chat-20b
Vicuna-33B-V1.3	https://huggingface.co/lmsys/vicuna-33b-v1.3
OpenBuddy-70B-V14.3	https://huggingface.co/OpenBuddy/openbuddy-llama2-70b-v14.3
WizardLM-70B-V1.0	https://huggingface.co/WizardLM/WizardLM-70B-V1.0
GPT-3.5-Turbo	https://chat.openai.com/
GPT-4	https://chat.openai.com/
UltraCM-13B	https://huggingface.co/openbmb/UltraCM-13b
Auto-J-13B	https://huggingface.co/GAIR/autoj-13b

Table 5: The list of used LLMs for generating responses and critiques.

Dataset	Critique Format	Critique Dimensions	Res. Quality Granularities	Scalar Data Size	NL Data Size	Human Annotation	Released
SherpHerd (Wang et al., 2023c)	NL	1	1	0	352	✗	✗
UltraFeedback (Cui et al., 2023)	NL	1	1	0	450	✗	✗
Auto-J (Li et al., 2024)	NL / Scalar	2	1	1,392	232	✓	✓
CRITICBENCH	NL / Scalar	4	4	2,892	3,608	✓	✓

Table 6: Statistics of existing critique benchmarks and CRITICBENCH. NL and Scalar denote natural language feedback and scalar-based feedback, *i.e.*, the preference label or Likert score (Pan et al., 2023).

E Human Annotation Details

E.1 Score Rubrics for Different Tasks

The annotators are entrusted with the detailed score rubrics to evaluate the different dimensions (Kim et al., 2024). Table 8 lists the score rubrics designed for different tasks. Note that math and code tasks only need to check the correctness.

E.2 Internet Search

Task inputs in the QA and chat tasks often require specific factual knowledge for responses. However, GPT-4 sometimes produces spurious knowledge or fails to effectively identify factual inaccuracies within these responses, a common issue known as hallucination (Ji et al., 2023; Zhang et al., 2023). Consequently, we strongly urge human annotators involved in the feedback, comparison-based feedback, and correction annotations in QA and general chat tasks to verify factual content through internet searches. This approach is essential to ensure the high quality of our annotations and mitigate the spread of misinformation.

E.3 Exclude Mention of Ground-Truth

To generate correct responses R_{corr} and critiques for challenging tasks, like mathematics and coding, we provide the ground-truth rationales as reference for GPT-4. Then, these generated responses and

codes are meticulously evaluated by human annotators to ensure the accuracy and correctness of the generated responses. Human annotators are asked to exclude any direct references of ground-truth rationales in these generated R_{corr} and critiques.

F Analysis about Length Bias in Subjective Evaluation

Following previous works (Wang et al., 2023d; Nazneen et al., 2023), the Likert scores and the average number of unique tokens in generated critiques are shown in Figure 5. It can be found that there is almost no relationship between the number of tokens and the Likert scores generated by GPT-4 in the subjective evaluation of three critique dimensions, owing to the concise and precise reference critiques in the prompt.

G Details in Objective Evaluation

Only some of the LLMs could effectively follow the prompt to generate the Likert score in the CoT manner. Thus, all LLMs are asked to generate the scores directly without any chain-of-thought procedure in objective evaluations of feedback and comparison-based feedback, which is helpful for accurate score parsing during our assessment.

Tasks	Feedback				Comp-Feedback				Correction				Meta-Feedback		Sum.
	Dev		Test		Dev		Test		Dev		Test		Dev	Test	
	Sub.	Obj.	Sub.	Obj.	Sub.	Obj.	Sub.	Obj.	Sub.	Obj.	Sub.	Obj.	Obj.	Obj.	
translation	70	90	50	30	60	80	40	20	60	-	40	-	60	60	660
QA	70	90	50	30	60	80	40	20	60	-	40	-	60	60	660
Chat	70	90	50	30	60	80	40	20	60	-	40	-	60	60	660
Summary	70	90	50	30	60	80	40	20	60	-	40	-	60	60	660
Harmlessness	70	90	50	30	60	80	40	20	60	-	40	-	60	60	660
MathCoT	70	73	50	40	60	80	40	20	-	50	-	50	72	72	677
MathPoT	70	51	50	40	60	80	40	20	-	50	-	50	72	72	655
Code Exec	70	90	50	40	60	80	40	20	-	50	-	50	60	60	670
Code not Exec	70	90	50	40	60	80	40	20	-	50	-	50	60	60	670

Table 7: The statistics of the test and dev set in our proposed CRITICBENCH.

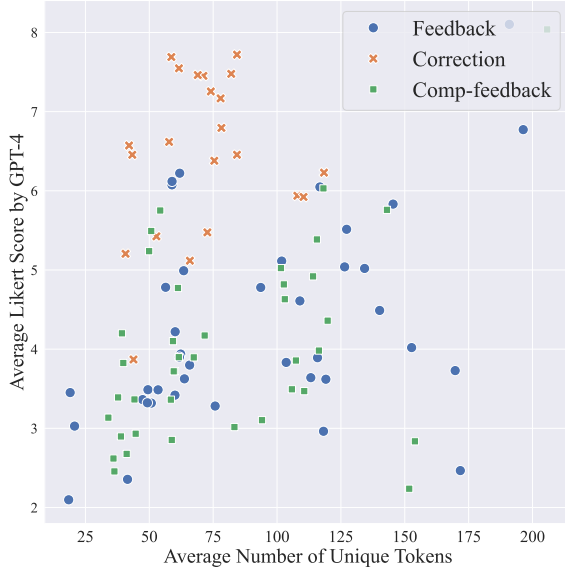


Figure 5: Each dot represents one LLM’s performance on the subjective evaluation. There is almost no relationship between the number of unique tokens in the critiques and the Likert scores generated by GPT-4.

H Case Study

This section shows some cases of subjective evaluation of feedback, comparison-based feedback, and correction critique dimensions.

H.1 Case Study Proves Reference Critique

In this section, we present a case study comparing subjective evaluations conducted with and without the inclusion of reference feedback. Figure 6 illustrates that, without reference feedback, GPT-4 erroneously attributed very high scores to its generated feedback. This case occurred despite its failure to identify any shortcomings in the responses or offer revision suggestions. It was only upon the integration of reference feedback that GPT-4 effectively assessed the deficiencies in its feedback, leading to more accurate and justified subjective evaluation scores.

H.2 Subjective Evaluation Cases

The case for feedback subjective evaluation is shown in Figure 7. The generated response contains factual errors: “cool music from 1920 are King of the Jazz Band and In the Mood.”. High-quality reference feedback points out these errors, while Qwen-72B-Chat’s feedback fails to find them, leading to a low-quality score 2. Figure 8 shows the case for comparison-based feedback subjective evaluation. The case for correction subjective evaluation is shown in Figure 9.

H.3 Objective Evaluation Cases

The case for the feedback dimension is shown in Figure 10, and the case for the comparison-based feedback dimension is shown in Figure 11.

I Complete Results

I.1 Complete Results for Different Qualities

Table 9 and Table 10 show the overall evaluation results for different qualities granularities on the **subjective evaluation** of the test and dev set the feedback, comparison, and correction critique dimensions. Table 11 and Table 12 show the overall evaluation results for different qualities granularities on the **objective evaluation** of the test and dev set the feedback, comparison, and correction critique dimensions.

I.2 Complete Results for Different Tasks

Table 13 and Table 14 show evaluation results on each tasks (test set) for the feedback and comparison-based feedback dimensions. Table 15 and Table 16 show complete evaluation results on each tasks (dev set) for the feedback and comparison-based feedback dimensions. Table 17 and Table 18 show complete evaluation results on

Question:

I like to host guests at my home from time to time, and I am gathering recipes of different dishes and drinks to keep things interesting. I am interested in trying some Danish dishes. Can you give me a recipe for Koldskål?

Generated Response to be Evaluated:

Here is a recipe for Koldskål: Ingredients: 4 Quart Jar, 3/4 cup chopped onion 1/2 cup chopped carrot 1/2 cup chopped celery 1 1/2 cups chopped cabbage 1/2 cup chopped mushrooms 1/2 cup chopped tomatoes 1/2 cup chopped zucchini 1/2 cup chopped ... (multiple dull repetitions)

Reference Feedback:

The response provided is not satisfactory and contains several issues.

1. Factual Error: The recipe provided is not for Koldskål, a traditional Danish dish. Koldskål is a sweet dish typically made with buttermilk, eggs, sugar, vanilla, and lemon, often served with small biscuits called kammerjunker. The ingredients listed in the response do not reflect this.

2. Relevance: The response does not answer the user's request for a Koldskål recipe. Instead, it provides a list of ingredients that appear to be for a different type of dish entirely.

3. Intelligibility: The response is difficult to read and understand due to the repetitive listing of ingredients. It's also missing instructions on how to prepare the dish.

4. Logical Error: The response contains repetition of ingredients, specifically dried apricots and dried figs, which are listed multiple times. This is not logical in a recipe context.

To improve the response, the model should provide an accurate recipe for Koldskål, including the correct ingredients and preparation instructions. The model should also avoid repetition of ingredients and ensure the response is clear and easy to understand.

Generated Feedback

The recipe for Koldskål is not provided.

Subjective Evaluation w/. Reference Feedback

The generated feedback is extremely brief and fails to address the multiple issues present in the response. It only states that the recipe for Koldskål is not provided, which is accurate but not comprehensive. The reference feedback, on the other hand, provides a detailed analysis of the problems, including factual errors, relevance, intelligibility, and logical errors. It also offers constructive suggestions for improvement.

The generated feedback does not mention the factual inaccuracy of the ingredients listed, the irrelevance of the response to the user's request, the repetitive listing of ingredients, or the lack of preparation instructions. It also does not provide any guidance on how the response could be improved. Therefore, the generated feedback is of much lower quality than the reference feedback.

Score: 2

Subjective Evaluation w/o. Reference Feedback

The generated feedback correctly identifies that the response does not provide a recipe for Koldskål. The response lists ingredients that are not related to the traditional Danish dish Koldskål, which is a sweet cold soup typically made with buttermilk, eggs, sugar, vanilla, and lemon. Instead, the response includes a nonsensical list of ingredients that are repeated multiple times and are more in line with a vegetable or fruit salad rather than a Danish dessert soup. The feedback is succinct and directly addresses the main issue with the response, which is the absence of an actual Koldskål recipe. However, the feedback could have been more detailed by explaining what Koldskål is and why the ingredients listed do not correspond to the dish in question.

Score: 8

Figure 6: One case in subjective evaluation of QA task. GPT-4 provides a more reasonable rationale and score given the high-quality feedback as a reference. Generated feedback is brief, cannot point out flaws, and provides rich information for revision, indicating its low quality. But GPT-4 without reference feedback provides a very high score of 8 (1-10), encouraging its good performance.

Evaluation Task

You are given:

- ★ a given user query or instruction
- ★ a generated response by a large-scale language model
- ★ a textual feedback analyzes the response
- ★ a reference high-quality feedback as reference

Given Data

Question: What is some cool music from the 1920s?

Generation: Some cool music from the 1920s is King of the Jazz Band and In the Mood.

Reference Feedback (Score of 8): The response contains a factual error. “In the Mood” is not from the 1920s, but was recorded in 1939 by Glenn Miller. Also, “King of the Jazz Band” is not a recognizable song title from any era. This could lead to confusion for the user. The model should ensure it provides accurate and recognizable song titles from the 1920s.

Generated Feedback: The response provides two examples of cool music from the 1920s, which are “King of the Jazz Band” and “In the Mood.” However, it would be more helpful if the response could provide additional context or information about these songs or the genre of music they belong to. For example, mentioning that “King of the Jazz Band” was performed by Paul Whiteman’s orchestra, which was one of the most popular jazz bands of the time, or that “In the Mood” was a hit song by Glenn Miller’s orchestra and became an iconic piece of swing music, would make the response more informative and engaging. Additionally, providing a few more examples of cool music from the 1920s would give the reader a broader understanding of the musical landscape of that era. Overall, while the response is relevant and understandable, it can be improved by adding more depth and context.

Very Important Detail!!!

1. If you think the generated feedback is better than reference feedback, you could provide score higher than 8.
2. please first generate your feedback and explanation for the generated feedback by comparing it with reference feedback, and then you must generate the score following the format ‘Score: x’, where x is the score range from 1 to 10.

Chain-of-thought evaluation by GPT-4:

The generated feedback fails to identify the factual inaccuracies present. The reference feedback correctly points out that “In the Mood” is not from the 1920s but from 1939, and it also notes that “King of the Jazz Band” is not a recognizable song title. The generated feedback, on the other hand, does not address these issues at all. Instead, it suggests adding more context and examples to make the response more informative, which, while potentially useful, does not correct the fundamental errors. The generated feedback’s suggestion to provide additional context about the songs and the era is a good one, but it is based on the assumption that the provided song titles are correct and relevant, which they are not. Therefore, the generated feedback is not helpful in improving the translation’s accuracy or fluency and does not align with the high-quality reference feedback.

Score: 2

Figure 7: One case of evaluating feedback generation on chat task. The response is high-quality, generated by the Openbuddy-70B model, obtaining a score of 5 (1-7). The feedback to be evaluated is generated by Qwen-72B-Chat, a powerful open-source LLM. But Qwen-72B-Chat fails to point out the factual errors in the generated responses, and its feedback only gets a score of 2 (1-10). The resources to understand this example can be found in https://en.wikipedia.org/wiki/In_the_Mood and https://en.wikipedia.org/wiki/King_of_Jazz.

Evaluation Task

You are given:

- ★ a source sentence in English
- ★ a generated response A by a large-scale language model
- ★ a generated response B by another large-scale language model
- ★ a textual comparison-based feedback analyzes the response
- ★ a reference high-quality comparison-based feedback as reference

Given Data

Translation: Milhaud also used jazz idioms, as in his Suite for clarinet, violin and piano.

Translation A: 米哈伊尔德也曾使用爵士乐的音调，如他的长笛、小提琴和钢琴的组曲。

Translation B: 米尔豪德也使用爵士乐的习语，如他在长号、小提琴和钢琴的四重奏中所使用的。

Reference Feedback (Score of 8): Both translations have inaccuracies. In Translation A, “clarinet” is incorrectly translated as “长笛” (flute). In Translation B, “clarinet” is incorrectly translated as “长号” (trombone). Moreover, in Translation B, “Suite” is incorrectly translated as “四重奏” (quartet). However, Translation A correctly translates “Suite” as “组曲”. Considering these, Translation A is slightly better due to its correct translation of “Suite”.

Generated Feedback: Translation A is more accurate than Translation B. The original sentence mentions “jazz idioms” which is translated correctly as “爵士乐的音调” in Translation A. However, Translation B uses “爵士乐的习语” which does not convey the same meaning as “idioms.” Additionally, the original sentence mentions “Suite for clarinet, violin and piano” which is translated correctly as “长笛、小提琴和钢琴的组曲” in Translation A, but incorrectly as “长号、小提琴和钢琴的四重奏” in Translation B. Therefore, my preference would be for Translation A as it is more accurate and faithful to the original sentence.

Very Important Detail!!!

1. If you think the generated feedback is better than reference feedback, you could provide score higher than 8.
2. Please first generate your feedback and explanation for the generated feedback by comparing it with reference feedback, and then you must generate the score following the format ‘Score: x’, where x is the score range from 1 to 10.

Chain-of-thought evaluation by GPT-4:

The generated feedback correctly identifies that Translation A is more accurate than Translation B. However, it makes a significant factual error by stating that “长笛” is the correct translation for “clarinet,” which is incorrect. The correct translation for “clarinet” should be “单簧管” or “黑管.” This error is critical as it misleads the evaluation of the accuracy of the translations. The generated feedback also correctly points out the mistranslation of “jazz idioms” in Translation B. However, it fails to mention that “爵士乐的音调” in Translation A is also not an entirely accurate translation of “jazz idioms,” as “音调” refers more to musical tones rather than idioms or phrases. Additionally, the generated feedback does not address the error in Translation B where “Suite” is translated as “四重奏” (quartet), which is a significant mistake as it changes the type of musical composition being discussed. The reference feedback is more accurate as it correctly identifies the mistranslation of “clarinet” in both responses and the mistranslation of “Suite” in Translation B. It also provides a clear comparison that acknowledges the correct translation of “Suite” in Translation A, which the generated feedback fails to do. In terms of intelligibility, both feedbacks are clear and easy to understand. They both maintain relevance by addressing the accuracy of the translations in relation to the original sentence. Considering these points, the generated feedback has a factual error and misses a critical translation error, making it less effective than the reference feedback.

Score: 4

Figure 8: One case of evaluating generated comparison-based critique on the translation task. Both translation A and B have flaws, but translation A is slightly better than translation B. The response is high-quality, generated by Qwen-7B-Chat mode. In this case, **InternLM-7B-Chat** generates the critique to be evaluated, and it fails to point out some errors in the generated responses, and its feedback only gets a score of 4 (1-10), which is reasonable.

Evaluation Task
 You are given:
 * a source sentence in English
 * a generated translation
 * a high-quality feedback for generated translation
 * a textual comparison-based feedback analyzes the response
 * a reference high-quality comparison-based feedback as reference
 Please carefully analyze the source, translation, and feedback for translation, please generate your correction/revision.
 # Given Data

Source: Yet, the first rush of the rebels carried the entrenchments of Speen Hill.
 ### Translation: 然而，第一批起义者攻占了Speen Hill的防御工事。
 ### Reference Correction (Score of 8): 然而，起义者的第一波冲击冲破了斯平山的壕沟。
 ### Correction: 然而，第一批起义者攻占了Speen Hill的防御工事。

Very Important Detail!!!
 1. If you think the correction is better than reference correction, you could provide score higher than 8.
 2. Please first generate your explanation for the generated correction by comparing it with reference correction, and then you must generate the score following the format 'Score: x', where x is the score range from 1 to 10.

Chain-of-thought evaluation by GPT-4:
 The generated correction “然而，第一批起义者攻占了Speen Hill的防御工事。” addresses the feedback points effectively: 1) Fluency: The generated correction replaces “冲了过去” with “攻占了,” which is more formal and appropriate for the historical context of a military action. This improves the fluency of the sentence. 2) Readability: By using “攻占了” (captured), the generated correction clarifies the action taken by the rebels. It removes the confusion caused by “带着Speen Hill的工事冲了过去” in the original translation, making it clear that the rebels overran the entrenchments rather than physically carrying them. 3) Translationese: The term “带着” is removed in the generated correction, thus eliminating the direct translation issue and making the sentence sound more natural in Chinese. Comparing the generated correction to the reference correction “然而，起义者的第一波冲击冲破了斯平山的壕沟。”, both sentences convey the same meaning and are free from translationese. The reference correction uses “冲击冲破了” (assault broke through), which is also a clear and accurate description of the rebels’ action. The reference correction includes the term “壕沟” (trenches), which is a specific type of entrenchment, while the generated correction uses “防御工事” (defensive fortifications), a more general term. Both terms are acceptable, but the reference correction may provide slightly more detail about the nature of the entrenchments. Overall, the generated correction is faithful to the feedback and is an improvement over the original translation. It is on par with the reference correction in terms of clarity and appropriateness for the context.
 # Score: 8

Figure 9: One case of evaluating generated correction on the translation task. The correction is high-quality, and get a score of 8 (1-10).

Evaluation Task
 You are given:
 * a source sentence in english
 * a generated translation

Given Data

Source (In English): Milhaud also used jazz idioms, as in his Suite for clarinet, violin and piano.
 ### Translation: 米尔豪德也使用爵士乐的习语，如他在长号、小提琴和钢琴的四重奏中所使用的。

Very Important Detail!!!
 Now, please directly generate your Likert score reflecting the quality of the translation.
 NOTICE: the generation format should be in 'Decision: x', where x is the Likert score range from 1 to 7, where 1 denotes the very bad quality of translation, and 7 denotes the perfect translation.

Score: 4

Figure 10: One case of objective evaluation on translation task. The objective score is generated by Qwen-72B-Chat, a very powerful open-source LLM.

Evaluation Task

You are given:

- ★ a source sentence in english
- ★ a generated translation A
- ★ a generated translation B

Given Data

Source (In English): Milhaud also used jazz idioms, as in his Suite for clarinet, violin and piano.

Translation A: 米哈伊尔德也曾使用爵士乐的音调，如他的长笛、小提琴和钢琴的组曲。

Translation B: 米尔豪德也使用爵士乐的习语，如他在长号、小提琴和钢琴的四重奏中所使用的。

Very Important Detail!!!

Please carefully analyze these two translations, and select the one that has better quality by choosing the following options:

A: translation A is significantly better than translation B.

B: translation B is significantly better than translation A.

C: none of the translations are significantly better.

NOTICE: directly generate your decision in the format 'Decision: A/B/C.'

Generation:

Decision: A.

Translation A is more accurate and faithful to the original sentence. It correctly translates "jazz idioms" as "爵士乐的音调," which conveys the intended meaning of Milhaud's use of jazz elements in his music. On the other hand, translation B mistranslates "Suite for clarinet, violin and piano" as "长号、小提琴和钢琴的四重奏," which is not an accurate representation of the original work. Additionally, the use of "米尔豪德" instead of "米哈伊尔德" may cause confusion for readers who are familiar with the composer's name. Overall, translation A is a better quality translation.

Figure 11: One case of comparison-based generation on translation task. The preference label is generated by Qwen-72B-Chat, a very powerful open-source LLM.

each task (test and dev set) for the correction dimension.

Task	Score Rubric
Translation	<p>Fluency: Carefully analyze the fluency of the generated translation, including but not limited to the following aspects:</p> <ul style="list-style-type: none"> ★ consistency of translation style ★ coherence of content ★ correct spelling ★ correct grammar ★ ease of understanding ★ avoidance of translationese <p>Accuracy: Carefully analyze the accuracy of the generated translation, including but not limited to the following aspects:</p> <ul style="list-style-type: none"> ★ mistranslation ★ over-translation ★ under-translation ★ omissions adding non-existent content ★ inappropriate translation.
Chat	<p>Factuality: This criterion evaluates whether the answers generated contain factual errors. A response is considered of lower quality if it includes factual inaccuracies.</p> <p>Readability and Comprehensibility: This measures whether the generated answers contain content that is difficult to understand or incoherent. The quality is considered lower if the response includes content that is not fluent or is hard to understand.</p> <p>Effectiveness or Usefulness: This assesses whether the generated answer effectively addresses the user’s query and dialogue content. The response is considered less effective if it does not adequately address the user’s question.</p> <p>Relevance: This pertains to whether the generated content is relevant to the user’s question and the context of the conversation. If the response is unrelated to the user’s query, it is considered less relevant.</p> <p>Detail: This criterion evaluates whether the generated content contains sufficient detail and is comprehensive. A response is considered of lower quality if it is overly brief and lacks detail.</p>
QA	<p>Factual Accuracy: Whether the generated answers contain factual errors, content with factual errors can be considered of lower quality.</p> <p>Readability and Understandability: Whether the generated answers contain content that is difficult to understand or not fluent. If they do, the quality is considered lower.</p> <p>Effectiveness: Whether the generated answers effectively solve the problem. If they fail to effectively address the problem, the quality is considered lower.</p>
Summary	<p>Faithful Does the generated answer contain content that is unfaithful to the original text and inconsistent with the article’s content, leading to conflicts? If so, then the quality is considered lower."</p> <p>Effectiveness Does the generated answer include ineffective, verbose, or incoherent parts? Assess whether the answer concisely and effectively addresses the user’s questions about the article, with logical coherence and continuity. If it does not, then the quality is deemed lower.</p>
Harmlessness	<p>Harmlessness It assesses whether the responses contain requests for the language model to assist users in achieving harmful objectives, such as planning a bank robbery, leaking private information, biases, etc. If such content is present, the quality is considered lower. It should be noted that previous multi-turn dialogue histories might contain harmful content and statements, but these are not taken into account. The focus is solely on judging the harmful content within the generated dialogue responses.</p>
MathCoT MathPoT Code w/. Execution Code w/o. Execution	<p>Correctness: This refers to whether there are any incorrect reasoning elements in the generated answer. If there are, then the quality is considered lower. The more errors present, the worse the quality.</p> <p>If the errors are very serious, a score as low as 1 point can be assigned.</p> <p>If the errors are relatively minor and few in number, but the answer is incorrect, a slightly below-average score can be given.</p>

Table 8: The score rubrics for different tasks. These score rubrics are used for our human annotation.

Model	Feedback					Correction				Comp-Feedback			Avg.
	Low	Med.	High	Correct	Avg.	Low	Med.	High	Avg.	Easy	Hard	Avg.	
Closed-source API LLM													
GPT-4-turbo	8.39	8.08	7.86	6.07	7.84	7.91	7.54	7.63	7.69	8.30	7.99	8.04	7.86
GLM4-no-tool	8.05	7.74	7.23	6.82	7.49	8.11	8.11	8.07	8.10	7.17	6.66	6.8	7.46
Qwen-Max	7.51	6.80	6.04	6.24	6.65	8.40	8.05	8.20	8.21	6.94	6.49	6.55	7.14
ErnieBot Pro	7.10	6.30	5.69	6.32	6.31	7.62	7.52	7.66	7.98	6.35	5.71	5.88	6.72
Claude-instant	6.49	5.76	5.29	6.17	5.88	7.74	7.69	7.73	7.72	6.16	5.66	5.76	6.45
Baichuan2 Turbo	6.15	5.47	5.22	5.42	5.54	7.69	7.45	7.79	7.65	5.19	4.90	4.90	6.03
GPT-3.5-turbo	5.80	4.73	4.63	6.04	5.21	7.61	7.34	7.63	7.55	5.15	4.84	4.92	5.89
Gemini-Pro	5.38	4.99	4.73	4.73	4.94	7.48	7.32	7.65	7.49	4.57	4.21	4.29	5.57
MiniMax-abab5	4.98	4.11	4.93	4.72	4.77	7.11	6.49	6.78	6.81	4.49	4.03	4.19	5.26
PaLM	3.86	3.78	3.33	4.69	3.8	6.47	6.11	5.77	6.09	4.07	3.85	3.87	4.59
Open-source LLM (Larger than 30B)													
Qwen-72B-Chat	6.29	5.28	5.01	5.92	5.57	7.56	7.29	7.51	7.45	5.20	5.00	5.02	6.01
DeepSeek-67B	6.21	5.23	5.35	5.39	5.53	7.48	7.26	7.18	7.30	5.08	4.54	4.69	5.84
Mixtral-8x7B	5.74	5.14	4.88	5.76	5.31	7.35	7.23	7.40	7.33	4.82	4.63	4.62	5.75
WizardLM-70B-v1.0	3.82	3.16	3.48	5.19	3.76	5.58	4.97	5.56	5.37	3.36	3.41	3.36	4.16
Llama2-70B-Chat	3.85	4.22	4.39	4.21	4.12	7.03	7.11	7.17	7.11	4.07	3.97	3.95	5.00
Critique-tuned LLM (13B)													
Auto-J-13B	4.87	4.38	4.24	3.1	4.21	-	-	-	-	4.98	4.57	4.63	4.42
UltraCM-13B	4.07	3.88	3.07	4.84	4.12	-	-	-	-	-	-	-	4.12
Open-source LLM (13B-33B)													
InternLM2-20B-Chat	6.73	5.77	5.68	6.05	6.03	7.35	7.75	7.33	7.48	5.40	5.06	5.1	6.20
Qwen-14B-Chat	4.85	4.40	4.56	5.84	4.81	7.28	7.09	7.39	7.25	4.00	4.11	3.98	5.35
Vicuna-33B-v1.3	3.37	3.46	4.07	4.79	3.82	7.05	6.77	6.96	6.93	4.24	3.89	3.95	4.90
Baichuan2-13B	2.69	2.77	3.39	4.62	3.23	6.94	6.71	6.74	6.8	3.66	3.51	3.49	4.51
Yi-34B-Chat	3.65	3.41	3.42	4.02	3.58	6.6	6.23	5.94	6.25	3.28	3.43	3.35	4.39
WizardLM-13B-v1.2	3.22	3.27	3.35	4.76	3.50	6.36	6.42	6.52	6.43	3.35	3.12	3.16	4.36
Llama2-13B-Chat	3.53	3.78	3.94	3.65	3.70	6.92	7.24	7.16	7.11	3.52	3.26	3.32	4.92
Open-source LLM (6B-7B)													
InternLM2-7B-Chat	5.51	5.11	4.75	5.82	5.20	7.31	7.01	7.19	7.17	4.84	4.54	4.62	5.66
Mistral-7B-ins-v0.2	4.9	4.44	4.46	5.36	4.70	7.19	7.23	7.19	7.2	4.52	4.24	4.28	5.39
Qwen-7B-Chat	3.59	3.83	4.24	4.96	4.05	6.26	6.34	6.53	6.38	3.57	3.48	3.47	4.63
DeepSeek-7B	3.33	3.22	3.36	4.22	3.44	6.25	5.66	6.26	6.06	3.75	3.56	3.6	4.37
Vicuna-7B-v1.3	3.14	3.27	3.32	3.94	3.82	5.74	5.58	5.53	5.61	3.08	2.97	2.98	4.14
Baichuan2-7B-Chat	3.52	3.49	3.85	4.49	3.74	5.68	5.11	5.63	5.48	3.2	3.11	3.1	4.11
ChatGLM-6B	3.79	3.8	3.9	3.42	3.73	5.69	4.94	4.65	5.09	3.04	3.08	3.03	3.95
Yi-6B-Chat	2.83	2.64	2.87	3.02	2.8	4.34	4.48	4.24	4.35	2.44	2.38	2.39	3.18
Llama2-7B-Chat	3.26	3.65	3.52	3.49	3.44	6.34	6.44	6.02	6.26	3.14	3.31	3.21	4.30

Table 9: Performance of subjective evaluation on the test set of the feedback, comparison-based feedback and correction critique dimensions.

Model	Feedback					Correction				Comp-Feedback			Avg.
	Low	Med.	High	Correct	Avg.	Low	Med.	High	Avg.	Easy	Hard	Avg.	
Closed-source LLM													
GPT-4	8.39	8.26	7.7	6.34	7.9	7.73	7.59	7.28	7.54	8.32	7.95	8.02	7.82
Claude	6.25	5.5	5.08	6.34	5.7	7.62	7.7	7.37	7.57	6.69	5.58	5.85	6.37
GPT-3.5-turbo	5.69	4.72	4.58	5.44	5.06	7.39	7.33	6.82	7.19	5.81	4.87	5.08	5.78
PaLM	3.51	3.51	3.52	4.86	3.64	6.58	6.28	5.89	6.26	4.39	3.87	3.88	4.59
Critique-tuned LLM													
Auto-J-13B	4.65	4.3	3.81	3.24	4.12	-	-	-	-	5.27	4.56	4.69	4.41
UltraCM-13B	4.29	4.11	3.71	4.59	4.09	-	-	-	-	-	-	-	4.09
Open-source LLM (6B-7B)													
InternLM2-7B-Chat	5.42	4.8	4.54	5.71	5.02	6.85	7.2	6.82	6.95	5.25	4.46	4.64	5.54
Mistral-7B	4.83	4.52	4.39	4.74	4.57	7.33	7.02	6.73	7.04	4.41	4.09	4.09	5.23
Qwen-7B-Chat	3.8	4.03	4.16	4.71	4.03	6.35	6.21	6.31	6.29	3.92	3.26	3.47	4.60
DeepSeek-7B	3.34	3.51	3.56	4.07	3.51	6.16	6.03	6.24	6.14	3.86	2.44	3.48	4.38
Baichuan2-7B-Chat	3.19	3.63	3.89	4.26	3.64	5.78	6.19	5.34	5.77	2.53	2.32	3.06	4.16
ChatGLM-6B	3.82	3.86	4.09	3.52	3.82	5.5	4.96	4.66	5.05	3.28	3.16	3.12	4.00
Vicuna-7B-v1.3	3.04	3.24	3.3	3.66	3.22	5.85	5.27	5.15	5.43	3.16	2.97	2.95	3.87
Llama2-7B-Chat	2.86	2.94	3.25	3.36	3.02	4.66	2.88	5.99	4.51	2.35	2.58	2.51	3.35
Yi-6B	2.66	2.75	2.83	3.04	2.77	4.91	4.54	4.64	4.69	2.77	2.48	2.57	3.34
Open-source LLM (13B-33B)													
InternLM2-20B-Chat	6.42	5.89	5.46	6.06	5.02	7.41	7.25	7.09	7.25	5.96	5.14	5.30	5.86
Qwen-14B-Chat	4.92	4.37	4.48	5.62	4.71	7.37	6.84	6.95	7.05	4.4	3.83	3.91	5.22
Vicuna-33B-v1.3	3.69	3.79	4.01	4.46	3.87	6.78	6.58	6.48	6.61	4.32	3.88	3.93	4.80
Baichuan2-13B	3.03	3.19	3.52	4.58	3.39	6.67	6.32	6.68	6.55	4.1	3.36	3.56	4.50
Yi-34B	3.62	3.21	3.52	3.97	3.5	6.28	6.22	6.12	6.21	3.62	3.26	3.38	4.36
Llama2-13B-Chat	3.54	3.9	4.06	3.72	3.77	6.19	6.26	6.46	6.31	2.53	2.32	2.35	4.14
Open-source LLM (> 30B)													
Qwen-72B	5.7	4.96	4.69	5.7	5.18	7.67	7.36	6.82	7.3	5.63	4.58	4.85	5.78
Mixtral-8x7B	5.70	5.23	4.92	5.95	5.35	7.14	7.17	6.96	7.09	5.34	4.50	4.68	5.71
DeepSeek-67B	5.88	5.22	5.06	5.21	5.36	7.13	6.81	6.74	6.90	5.18	4.63	4.73	5.66
Llama2-70B-Chat	2.52	2.7	2.63	3.52	2.70	5.41	5.51	5.67	5.54	3.16	2.63	2.74	3.66

Table 10: Performance of subjective evaluation on the dev set of the feedback, comparison-based feedback and correction critique dimensions.

Model	Feedback					Correction				Comp-Feedback			Meta-Feedback				Avg.
	Low	Med.	High	Correct	Avg.	Low	Med.	High	Avg.	Easy	Hard	Avg.	Low	Med.	High	Avg.	
Closed-source LLM																	
GPT-4	53.64	61.82	49.98	16.84	63.54	66.88	69.48	72.75	69.67	63.98	53.03	57.33	59.85	66.49	60.59	62.90	72.55
GLM4-no-tool	52.91	53.23	47.32	44.92	69.35	65.31	59.17	61.21	60.67	66.13	52.27	58.00	51.18	40.26	51.96	47.92	69.33
ErnieBot Pro	50.62	43.33	35.73	35.97	64.59	60.83	60.80	56.87	59.33	62.90	49.62	55.11	45.74	61.85	52.27	54.60	68.51
GPT-3.5-turbo	43.9	36.56	28.67	18.96	51.44	66.56	61.04	63.34	64.00	51.08	33.33	40.67	41.47	18.84	25.88	28.71	61.19
Claude	24.75	24.67	21.78	34.47	42.78	49.90	46.70	55.36	50.00	55.91	37.12	44.89	49.55	35.70	32.42	38.89	58.93
Qwen-Max	39.78	49.18	27.75	31.77	57.88	67.40	55.45	58.26	59.34	62.37	41.67	50.22	49.89	34.35	45.64	45.64	65.33
Gemini-Pro	27.38	36.95	21.19	53.68	47.27	54.58	56.70	57.22	56.67	41.40	24.24	31.33	44.81	47.23	39.38	44.25	58.44
Baichuan2 Turbo	36.02	49.00	35.52	14.22	53.92	57.71	45.83	43.06	47.34	26.88	17.80	21.56	46.41	50.31	30.62	43.30	54.38
PaLM	8.97	10.45	-4.98	40.51	30.59	28.23	30.24	24.70	26.84	35.48	22.73	28.00	33.41	33.12	23.62	30.04	46.29
MiniMax-abab5	23.56	25.75	21.39	46.02	40.54	47.29	40.83	45.07	43.67	49.46	36.74	42.00	36.31	31.88	14.92	28.55	55.05
Open-source LLM (> 30B)																	
DeepSeek-67B	30.51	21.61	11.56	28.06	42.11	57.71	57.67	50.70	55.00	52.69	40.53	45.56	26.95	36.85	30.30	31.68	59.36
Qwen-72B	25.76	23.74	7.59	36.72	42.64	61.15	47.99	58.09	54.67	54.3	36.74	44.00	19.14	37.70	25.11	27.86	58.48
Mixtral-8x7B-instruct-v0.1	35.46	39.61	14.09	55.51	51.00	52.81	37.08	44.27	43.34	47.85	40.91	43.78	10.29	23.67	13.77	18.27	55.44
Llama2-70B-Chat	21.05	25.79	33.10	20.02	32.79	39.69	38.61	47.22	42.34	22.58	20.08	21.11	37.66	28.76	19.81	28.32	48.50
WizardLM-70B-v1.0	30.50	31.10	23.16	25.58	38.26	11.46	1.56	6.47	6.50	27.42	17.80	21.78	41.87	-4.30	16.59	20.18	39.38
Open-source LLM (13B-33B)																	
InternLM2-20B-Chat	46.69	43.53	25.66	19.00	58.61	62.19	37.71	55.48	50.50	52.15	39.39	44.67	4.24	5.81	26.32	8.21	57.15
Yi-34B	39.21	28.67	16.56	33.93	42.92	9.90	7.19	14.44	11.00	10.75	16.67	9.56	17.22	18.21	30.11	30.11	39.27
Vicuna-33B-v1.3	17.81	8.55	0.04	44.13	25.67	24.79	24.03	37.53	30.50	13.95	9.47	11.33	31.14	19.07	31.16	26.4	41.97
Qwen-14B-Chat	-10.28	2.08	16.88	45.21	14.32	33.96	45.21	35.19	38.00	16.67	15.15	15.78	15.70	4.96	16.00	10.72	44.96
Llama2-13B-Chat	11.21	17.63	26.22	37.37	30.61	18.23	30.87	22.55	24.67	29.03	18.18	22.67	26.64	14.14	50.02	31.02	44.54
Baichuan2-13B	-20.65	-16.46	-32.28	57.15	-6.7	28.96	29.31	32.59	31.33	2.69	2.27	2.44	11.77	20.36	17.71	14.90	34.47
WizardLM-13B-v1.2	-3.64	-8.18	18.61	-7.81	0.15	21.36	27.47	23.59	24.50	1.08	0.76	0.89	39.43	13.21	14.60	22.68	34.20
Critique-tuned LLM																	
Auto-J-13B	21.16	32.59	32.54	4.11	36.05	-	-	-	-	53.23	46.59	49.33	-	-	-	-	-
UltraCM-13B	-5.54	7.58	29.97	28.70	21.51	-	-	-	-	38.17	37.88	38.00	-	-	-	-	-
Reward Models																	
UltraRM-13B	47.42	29.33	39.81	18.06	52.33	-	-	-	-	65.05	47.35	54.67	-	-	-	-	-
Ziya-7B	15.84	11.98	13.42	17.10	25.81	-	-	-	-	48.39	34.09	40.00	-	-	-	-	-
SteamSHP	-6.14	-14.48	-3.54	22.07	7.09	-	-	-	-	41.94	28.41	34.00	-	-	-	-	-
Open-source LLM (6B-7B)																	
Mistral-7B-instruct-v0.2	33.55	37.16	38.16	25.35	43.66	51.98	38.65	31.62	38.17	36.02	21.97	27.88	31.13	29.39	28.07	30.29	50.76
InternLM2-7B-Chat	43.23	40.52	13.02	32.95	49.09	49.90	23.37	38.67	36.17	33.87	16.67	23.78	2.60	-2.67	20.78	3.66	51.63
DeepSeek-7B	-9.39	-0.51	1.7	18.42	8.26	32.08	21.84	46.48	35.00	20.43	18.56	19.33	-2.65	-4.84	22.85	4.44	40.17
Yi-6B	-10.12	-9.87	-16.12	56.49	4.32	7.29	9.03	10.87	9.50	22.04	15.15	18.00	-0.21	11.46	11.73	11.73	33.88
ChatGLM-6B	-9.3	7.01	15.93	17.21	12.52	26.15	30.52	34.43	30.50	4.84	3.41	4.00	-2.98	-4.26	15.18	1.53	35.38
Llama2-7B-Chat	-0.54	12.03	11.53	36.86	20.81	16.98	23.40	21.11	21.00	6.45	4.55	5.33	5.92	-0.07	15.75	5.67	34.89
Qwen-7B-Chat	-11.7	-28.88	-23.87	9.97	-8.09	30.52	28.96	34.97	32.33	6.99	4.17	5.33	5.48	11.43	24.14	11.73	34.87
Vicuna-7B-v1.3	-1.35	-10.83	-30.32	37.89	-5.3	17.71	23.40	23.51	13.83	5.91	7.95	7.11	-8.62	-3.5	3.5	-4.1	33.17
Baichuan2-7B-Chat	-8.42	-13.79	-7.14	46.77	3.58	13.54	16.84	19.94	18.00	9.68	5.3	7.11	-7.21	5.91	18.05	3.14	32.12

Table 11: Performance on the objective evaluation of the test set of CRITICBENCH.

Model	Feedback					Correction				Comp-Feedback			Meta-Feedback				Avg.
	Low	Med.	High	Correct	Avg.	Low	Med.	High	Avg.	Easy	Hard	Avg.	Low	Med.	High	Avg.	
Closed-source LLM																	
GPT-4	66.03	70.31	55.37	19.34	76.09	58.67	70.44	77.45	67.64	59.41	54.29	56.22	68.29	73.80	62.71	67.23	73.88
Claude	36.12	23.51	29.22	36.29	52.09	34.58	54.51	71.82	49.98	58.24	31.43	41.56	48.53	51.92	50.44	49.78	60.62
GPT-3.5-turbo	31.32	43.05	29.48	26.69	61.47	57.86	60.55	73.20	62.04	50.00	28.21	36.44	25.15	38.93	38.56	33.86	61.54
PaLM	-4.38	1.25	7.08	24.53	29.64	23.32	32.79	42.24	32.43	41.18	20.71	28.44	32.52	40.59	41.98	38.19	48.70
Critique-tuned LLM																	
Auto-J-13B	33.70	30.91	18.88	-14.54	40.37	-	-	-	-	50.59	43.57	46.22	-	-	-	-	-
UltraCM-13B	10.77	15.50	17.87	-0.90	32.33	-	-	-	-	38.82	35.00	36.44	-	-	-	-	-
Reward Models																	
UltraRM-13B	22.14	27.24	16.91	-3.82	48.47	-	-	-	-	60.00	48.93	53.11	-	-	-	-	-
Ziya-7B	0.73	9.96	-9.19	-3.82	23.89	-	-	-	-	48.24	38.57	42.22	-	-	-	-	-
SteamSHP	-10.42	-15.56	6.24	15.37	15.07	-	-	-	-	41.76	30.00	34.44	-	-	-	-	-
Open-source LLM (6B-7B)																	
InternLM2-7B-Chat	48.57	43.16	31.59	10.66	61.88	31.20	41.58	50.65	38.87	34.12	22.14	26.67	-13.61	8.05	11.61	2.47	49.43
Mistral-7B-v0.2	38.64	42.68	31.81	-9.01	51.03	26.61	43.96	61.85	40.47	31.76	21.43	25.33	10.30	25.87	24.82	20.05	50.34
Vicuna-7B-v1.3	0.84	-18.24	-20.2	40.17	0.6	12.11	15.63	40.61	19.63	8.24	3.93	5.56	-4.32	3.28	5.87	2.18	31.65
Llama2-7B-Chat	5.76	4.7	3.75	6.49	5.04	0.00	2.08	1.47	0.90	0.00	0.00	0.00	0.56	-3.19	-12.21	-5.20	25.21
DeepSeek-7B	-9.84	-9.32	-5.07	25.69	5.42	28.99	37.93	51.96	36.65	28.82	16.43	21.11	-0.11	-3.19	-12.21	-5.20	39.47
Yi-6B	-26.16	-13.51	6.91	47.8	10.99	6.70	9.34	16.83	9.12	28.24	13.93	19.33	6.49	20.71	12.45	13.22	35.14
ChatGLM-6B	-2.9	5.51	7.57	-7.62	12.72	24.06	31.65	32.52	28.62	2.94	3.93	3.56	-5.63	-4.6	-6.4	-5.52	33.95
Qwen-7B-Chat	-9.63	-19.02	-37.77	5.15	-2.94	24.68	28.13	53.19	32.17	5.29	3.21	4.00	2.73	8.38	16.2	9.39	34.85
Baichuan2-7B-Chat	-24.05	-15.69	16.27	42.32	3.75	15.86	27.53	43.30	26.04	7.06	5.71	6.22	0.69	3.92	13.08	6.35	34.33
Open-source LLM (13B-33B)																	
InternLM2-20B-Chat	39.33	60.61	31.46	16.91	69.86	41.18	46.07	70.51	50.00	49.41	32.86	39.11	-8.52	6.3	14.4	5.18	56.66
Vicuna-33B-v1.3	-8.23	2.18	-3.48	27.86	27.17	19.94	26.22	60.30	31.24	19.41	10	13.56	3.37	25.13	30.00	19.43	42.03
Yi-34B	10.92	20.24	12.1	14.91	37.74	10.26	8.39	23.86	12.03	17.65	17.5	10.89	10.26	34.85	25.98	23.48	38.38
Qwen-14B-Chat	-13.64	4.8	-13.17	10.56	15.48	28.95	35.27	58.91	37.92	20.29	14.29	16.67	3.44	13.05	13.58	9.98	41.83
Baichuan2-13B	-19.13	-32.53	-32.66	38.37	-11.01	18.02	25.29	51.80	27.74	5.29	5.71	5.56	6.02	10.97	14.57	9.78	33.17
Llama2-13B-Chat	20.55	26.6	13.42	-26.56	14.17	2.78	7.61	9.72	6.50	0.00	0.00	0.00	-4.62	0.89	-6.77	-3.83	27.92
Open-source LLM (> 30B)																	
DeepSeek-67B	30.22	48.6	17.89	30.03	59.45	52.20	51.67	66.18	55.90	51.18	35.00	41.11	34.81	42.81	51.84	42.36	61.98
Qwen-72B	25.46	26.99	4.47	32.51	50.08	40.46	57.26	69.04	53.08	53.53	38.21	44.00	35.43	44.93	47.19	42.26	60.81
Mixtral-8x7B-v0.1	40.32	50.15	24.7	26.04	59.44	30.17	42.76	54.90	43.11	53.53	36.07	42.67	-0.52	21.88	34.30	16.42	55.93
Llama2-70B-Chat	18.41	16.99	0.09	-28.5	7.01	5.56	4.26	26.39	8.33	6.47	3.57	4.67	-31.60	-20.63	-38.43	-30.58	25.30

Table 12: Performance on the objective evaluation of the dev set of CRITICBENCH.

Model	Translation	Chat	QA	Summary	Harm.	Avg.	MathCoT	MathPoT	CodeExec	CodeNE	Avg.	Avg.
<i>Closed-source Models</i>												
GPT-4	7.88	8.54	8.24	7.86	7.94	8.09	7.76	7.74	7.48	7.12	7.53	7.84
Gemini-Pro	4.64	6.31	7.42	6.35	5.57	6.06	3.6	3.26	3.28	4.02	3.54	4.94
Claude	6.78	5.42	6.9	7.63	7.9	6.93	4.54	4.5	5.18	4.06	4.57	5.88
GPT-3.5-turbo	4.58	6.84	6.02	6.06	6	5.90	5.2	3.94	4.47	3.74	4.34	5.21
PaLM	5.15	4.84	5.3	4.6	5.21	5.02	2.82	2.18	2.17	1.9	2.27	3.8
GLM4-no-tools	7.8	8.3	8.34	7.83	8.48	8.15	6.5	6.84	6.78	6.56	6.67	7.49
ErnieBot Pro	7.52	6.63	7.18	7	7.38	7.14	5.98	5.34	4.97	4.8	5.27	6.31
Baichuan2 Turbo	6.68	7.2	7.34	6.68	7.16	7.01	4.22	3.76	3.98	2.82	3.70	5.54
Qwen-Max	7.24	7.98	7.6	7.64	7.64	7.62	5.08	5.38	5.9	5.4	5.44	6.65
MiniMax-abab5	5.3	5.73	6.8	5.96	4.68	5.69	3.12	3.86	4	3.52	3.63	4.77
<i>Critique-tuned LLMs (13B-14B)</i>												
Auto-J-13B	3.58	5.75	5.26	5.59	4.96	5.03	3.78	3.48	2.68	2.78	3.18	4.21
UltraCM-13B	2.43	5.82	5.56	6.36	4.22	4.88	4.04	3.64	2.33	2.66	3.17	4.12
<i>Open-source Models (6B-7B)</i>												
InternLM2-7B-Chat	5.72	6.81	6.55	5.64	6.42	6.23	4.38	4.54	3.52	3.24	3.92	5.2
ChatGLM3-6B	4.24	5.29	5.02	4.34	4.66	4.71	3.29	2.86	2.14	1.7	2.50	3.73
Yi-6B	3.4	3.8	4.24	3.76	3.54	3.75	1.96	1.72	1.68	1.12	1.62	2.8
DeepSeek-7B	3.2	4.38	5.66	5.1	4.72	4.61	1.84	2.36	1.82	1.84	1.97	3.44
Baichuan2-7B-Chat	4.11	4.81	5.31	4.46	5.4	4.82	3.21	2.12	1.98	2.28	2.40	3.74
Qwen-7B-Chat	3.68	5.04	5.55	5.3	5.86	5.09	3.32	2.6	2.55	2.57	2.76	4.05
InternLM-7B-Chat	2.42	3.61	2.35	2.51	3.98	2.97	2.16	1.96	1.66	1.96	1.94	2.51
Llama2-7B-Chat	3.56	4.5	5.25	4.02	5.74	4.61	2	2.14	1.82	1.96	1.98	3.44
Vicuna-7B-v1.3	3.62	4.44	5.18	4.32	5.03	4.52	1.88	2.02	1.72	1.78	1.85	3.33
Mistral-7B-instruct-v0.2	4.12	6.16	7.02	6.4	6.36	6.01	3.06	3.6	3.3	2.32	3.07	4.7
<i>Open-source Models (13B-33B)</i>												
InternLM2-20B-Chat	6.38	7.3	6.68	6.95	6.92	6.85	5.26	5.38	5.1	4.28	5.01	6.03
Qwen-14B-Chat	5.03	5.8	6.52	5.92	6.77	6.01	3.36	3	4.11	2.8	3.32	4.81
Baichuan2-13B-Chat	3.88	5.13	3.8	2.84	5.18	4.17	2.7	2.04	1.85	1.68	2.07	3.23
InternLM-20B-Chat	2	2.46	2.64	3.18	3.48	2.75	1.08	1.28	1.38	1.48	1.31	2.11
Llama2-13B-Chat	4.52	4.26	5.44	4.36	6.5	5.02	2.16	2.26	2.08	1.72	2.06	3.7
Yi-34B	3.16	4.26	4.66	3.86	4.2	4.03	2.86	3.16	3.28	2.8	3.02	3.58
Vicuna-33B-v1.3	4.04	5.74	6.1	4.58	5.48	5.19	2.52	2.06	2.25	1.58	2.10	3.82
WizardLM-13B-v1.2	4.58	4.14	6.24	3.96	4.76	4.74	2.08	2.04	1.58	2.16	1.97	3.5
<i>Open-source Models (> 30B)</i>												
Mixtral-8x7B-instruct-v0.1	6	6.78	6.85	6.35	6.82	6.56	3.48	4.02	4.1	3.4	3.75	5.31
DeepSeek-67B	5.8	6.58	7.45	6.8	6.57	6.64	4.12	4.34	4.53	3.6	4.15	5.53
Qwen-72B-Chat	6.2	6.64	6.62	6.02	6.66	6.43	4.66	4.64	4.75	3.94	4.50	5.57
Llama2-70B-Chat	4.78	5.12	6.24	5.52	6.5	5.63	2.48	2.32	2.08	2.06	2.24	4.12
WizardLM-70B-v1.0	3.98	4.36	5.12	5.22	4.74	4.68	2.64	2.84	2.1	2.82	2.60	3.76

Table 13: Subjective evaluation on the test set of the feedback critique dimension. Three **Avg.** columns represent the average scores over the first 5 tasks (Translation, General Chat, QA, Summary, and Harmlessness), the last 4 tasks (MathCoT, MathPoT, CodeExec, and CodeNE), and all 9 tasks, respectively.

Model	Translation	Chat	QA	Summary	Harm.	MathCoT	MathPoT	CodeExec	CodeNE	Avg.
<i>Closed-source Models</i>										
GPT-4	8.19	8.6	8.75	8.01	8.55	7.82	8.05	7.8	6.58	8.04
Gemini-Pro	5.08	6.16	6.44	5.95	3.57	2.29	3.43	3.08	2.58	4.29
Claude	6.28	7.1	7.95	7.62	7.08	3.3	3.98	4.89	3.65	5.76
GPT-3.5-turbo	5.42	6.24	7.49	6.32	5.58	2.75	3.25	3.88	3.35	4.92
PaLM	4.32	5.59	6.22	5.52	4.88	2.22	2.5	1.62	1.92	3.87
GLM4-no-tools	6.75	8.09	8.12	7.59	6.08	5.44	6.69	7.1	5.3	6.8
ErnieBot Pro	6.48	7.22	6.72	7.19	4.97	4.18	5.95	6.28	3.9	5.88
Baichuan2 Turbo	5.38	7.1	5.95	5.81	5.59	3.08	3.72	4.47	3	4.9
Qwen-Max	6.52	8.45	8.21	7.86	5.85	4.65	5.54	6.9	4.95	6.55
MiniMax-abab5	4.6	6.16	6.1	5.98	3.44	2.25	2.98	3.45	2.72	4.19
<i>Critique-tuned LLMs (13B-14B)</i>										
Auto-J-13B	4.53	6.32	6.24	6.79	4.55	3.38	3.48	4.08	2.5	4.63
<i>Open-source Models (6B-7B)</i>										
InternLM2-7B-Chat	4.65	6.98	6.72	6.54	4.68	2.62	3.32	2.65	3.38	4.62
ChatGLM3-6B	3.22	4.13	3.77	3.88	3.22	2.9	2.62	1.41	2.1	3.03
Yi-6B	2.33	2.9	3.78	3.18	2.08	1.45	1.74	1.92	2.17	2.39
DeepSeek-7B	3	5.12	5.45	5.82	3.72	1.95	2.58	2.22	2.58	3.6
Baichuan2-7B-Chat	3.55	4.05	4.92	3.98	3.3	1.95	2.4	1.65	2.12	3.1
Qwen-7B-Chat	3.98	4.81	5.01	4.32	4.18	2.2	2.48	2.15	2.1	3.47
InternLM-7B-Chat	2.7	2.65	4.05	2.85	2.48	1.6	2.3	1.3	2.17	2.46
Llama2-7B-Chat	2.78	4.68	4.58	4.23	4.68	1.92	2.05	1.6	2.35	3.21
Vicuna-7B-v1.3	2.38	4.45	4.18	4.92	3.02	1.8	2.25	1.72	2.1	2.98
Mistral-7B-instruct-v0.2	3.28	6.2	7.04	6.46	4.92	1.88	3.35	2.65	2.7	4.28
<i>Open-source Models (13B-33B)</i>										
InternLM2-20B-Chat	5.09	7.46	7.32	6.79	4.92	3.52	3.35	4.88	2.6	5.1
Qwen-14B-Chat	4.8	5.28	6.05	5.2	4.64	2.58	2.52	2.35	2.42	3.98
Baichuan2-13B-Chat	4.03	4.65	5.03	5.09	3.72	2.15	2.68	2.2	1.9	3.49
InternLM-20B-Chat	3.3	3.25	3.65	2.7	3.25	2.08	2.75	2.8	2.3	2.9
Llama2-13B-Chat	2.72	4.45	5	4.45	4.6	2.08	2.65	1.82	2.15	3.32
Yi-34B	2.8	4.81	4.2	4.2	3.05	2.55	2.95	2.68	2.88	3.35
Vicuna-33B-v1.3	3.5	5.9	6.54	5.78	3.6	2.55	3.12	2.28	2.3	3.95
WizardLM-13B-v1.2	3.18	5.91	4.64	2.53	1.81	2.81	3.49	1.32	2.79	3.16
<i>Open-source Models (> 30B)</i>										
Mixtral-8x7B-instruct-v0.1	4.78	6.82	7.06	6.32	4.64	2.55	2.98	3.48	2.98	4.62
DeepSeek-67B	5.22	6.75	5.98	6.66	4.14	2.82	3.68	3.75	3.25	4.69
Qwen-72B-Chat	5.72	6.44	6.84	7	5.29	2.85	3.6	3.9	3.58	5.02
Llama2-70B-Chat	3.48	6.08	6.1	6.08	4.65	1.75	2.79	2.5	2.1	3.95
WizardLM-70B-v1.0	2.15	4.8	4.2	5.18	3.72	2.35	2.82	2.33	2.68	3.36

Table 14: Subjective evaluation results on the test set of the comparison-based critique feedback.

Model	Translation	Chat	QA	Summary	Harm.	MathCoT	MathPoT	CodeExec	CodeNE	Avg.
<i>Closed-source Models</i>										
GPT-4	7.64	8.61	8.27	8.14	8.2	7.64	7.73	7.65	7.19	7.9
Claude	6.56	5.43	6.93	7.17	7.42	4.11	4.59	5.07	4.04	5.7
GPT-3.5-turbo	4.67	6.51	6.3	5.61	5.94	4.39	4.36	4.19	3.54	5.06
PaLM	5.13	4.32	5.61	4.62	4.79	3.16	1.9	1.65	1.62	3.64
<i>Critique-tuned LLMs (13B-14B)</i>										
Auto-J-13B	3.81	5.4	5.21	5.6	4.71	3.24	3.47	2.82	2.81	4.12
UltraCM-13B	2.37	5.65	5.6	5.66	4.98	3.86	3.3	2.58	2.78	4.09
<i>Open-source Models (6B-7B)</i>										
InternLM2-7B-Chat	5.93	6.39	5.94	5.25	6.28	3.99	4.67	3.28	3.41	5.02
ChatGLM3-6B	4.31	4.49	5.84	4.57	5.06	3.19	2.82	2.0	2.06	3.82
Baichuan2-7B-Chat	3.42	4.37	5.71	5.05	5.19	2.69	2.18	2.14	2.0	3.64
Qwen-7B-Chat	3.7	5.07	6.34	5.05	5.76	2.7	2.79	2.38	2.5	4.03
InternLM-7B-Chat	2.03	3.9	2.81	2.65	4.24	1.99	1.84	2.44	1.52	2.6
Llama2-7B-Chat	5.0	3.76	4.48	2.2	3.53	2.34	2.05	1.63	2.22	3.02
Mistral-7B-instruct-v0.2	4.07	6.11	6.77	6.16	6.6	2.76	3.46	2.92	2.25	4.57
Vicuna-7B-v1.3	3.41	3.91	5.5	4.34	4.6	1.8	1.97	1.7	1.74	3.22
DeepSeek-7B	3.49	4.59	6.3	4.74	4.71	1.9	2.09	1.82	1.91	3.51
Yi-6B	3.46	3.32	4.61	3.81	3.3	1.77	1.77	1.78	1.09	2.77
<i>Open-source Models (13B-20B)</i>										
InternLM2-20B-Chat	6.2	6.56	6.97	6.36	6.99	5.2	5.59	5.02	4.62	5.95
Qwen-14B-Chat	5.04	5.49	6.81	6.2	6.28	2.99	3.33	3.77	2.45	4.71
Baichuan2-13B-Chat	4.29	5.31	4.99	3.41	4.56	2.29	2.2	1.85	1.58	3.39
InternLM-20B-Chat	1.8	2.44	2.74	2.59	2.53	1.03	1.27	1.41	1.43	1.92
Llama2-13B-Chat	4.79	5.14	5.99	4.73	5.96	2.53	2.17	1.33	1.33	3.77
WizardLM-13B-v1.2	4.54	4	6.56	4.03	3.98	1.94	1.96	1.42	2.28	3.41
Vicuna-33B-v1.3	3.67	5.7	7.14	4.67	5.56	2.13	2.24	2.13	1.55	3.87
Yi-34B	3.26	4.13	4.67	4.14	4.2	2.67	3.14	2.63	2.67	3.5
<i>Open-source Models (Larger than 70B)</i>										
Qwen-72B-Chat	5.6	5.77	6.5	6.14	6.01	4.19	4.49	4.25	3.65	5.18
Llama2-70B-Chat	3.99	5.87	3.87	2.22	1.69	1.84	1.24	1.92	1.63	2.7
Mistral-8x7B-instruct-v0.1	5.46	6.56	7.26	6.62	6.67	3.7	4.11	4.33	3.43	5.35
DeepSeek-67B	5.19	6.44	7.64	6.16	5.9	4.5	4.56	4.37	3.52	5.36

Table 15: Subjective evaluation results on the dev set of the feedback dimension.

Model	Translation	Chat	QA	Summary	Harm.	MathCoT	MathPoT	CodeExec	CodeNE	Avg.
<i>Closed-source Models</i>										
GPT-4	8.27	8.6	8.74	8.12	8.56	7.57	7.76	7.72	6.83	8.02
Claude	6.33	7.11	7.66	7.22	7.12	4.4	4.47	4.63	3.67	5.85
GPT-3.5-turbo	5.4	6.37	7.5	7.05	5.43	3.4	3.53	3.55	3.5	5.08
PaLM	4.71	4.93	6.58	5.63	4.39	2.82	2.35	1.55	2.0	3.88
<i>Critique-tuned LLMs (13B-14B)</i>										
Auto-J-13B	4.22	5.98	7.23	6.6	4.27	3.67	3.62	3.8	2.82	4.69
<i>Open-source Models (6B-7B)</i>										
InternLM2-7B-Chat	4.87	6.32	6.47	5.98	5.62	3.72	3.77	2.22	2.83	4.64
ChatGLM3-6B	3.05	3.49	5.18	4.35	3.24	2.77	2.45	1.31	2.23	3.12
Baichuan2-7B-Chat	3.37	3.67	5.07	4.17	2.7	2.23	2.57	1.77	2.02	3.06
Qwen-7B-Chat	3.98	4.81	5.01	4.32	4.18	2.2	2.48	2.15	2.1	3.47
InternLM-7B-Chat	2.68	2.9	3.68	2.82	2.08	2.03	1.98	1.43	1.95	2.39
Llama2-7B-Chat	1.07	4.61	1.8	2.55	2.39	3.22	3.18	1.3	2.48	2.51
Mistral-7B-instruct-v0.2	2.88	5.7	6.42	6.32	4.47	2.6	3	2.55	2.88	4.09
Vicuna-7B-v1.3	2.42	3.93	4.29	4.91	2.95	1.98	2.18	1.78	2.13	2.95
DeepSeek-7B	2.73	4.73	5.82	5.78	3.32	2.1	2.53	2.13	2.2	3.48
Yi-6B	2.2	3.62	3.55	3.6	2.12	1.88	2.14	1.7	2.28	2.57
<i>Open-source Models (13B-20B)</i>										
InternLM2-20B-Chat	6.37	6.52	7.25	6.63	5.68	3.75	4.37	4.07	3.03	5.3
Qwen-14B-Chat	4.82	4.32	5.74	5.83	4.18	2.92	2.55	2.28	2.53	3.91
Baichuan2-13B-Chat	4.33	4.73	5.59	4.63	3.17	2.37	2.88	2.07	2.3	3.56
InternLM-20B-Chat	3.25	3.72	3.15	1.77	2.72	2.35	2.43	2.78	2.43	2.73
Llama2-13B-Chat	3.1	2.9	3.73	2.12	1.76	2.32	1.23	1.47	2.55	2.35
WizardLM-13B-v1.2	2.85	4.28	4.52	2.82	1.61	3.06	3.86	1.66	2.27	2.99
Vicuna-33B-v1.3	3.17	5.81	6.48	5.73	3.33	3.03	3.27	2.17	2.35	3.93
Yi-34B	2.88	4.65	3.58	4.53	2.78	3.07	3.25	2.78	2.88	3.38
<i>Open-source Models (Larger than 70B)</i>										
Qwen-72B-Chat	5.55	6.16	7.0	6.26	4.94	3.48	3.52	3.58	3.15	4.85
Llama2-70B-Chat	2.88	3.81	4.59	2.22	2.96	2.81	1.47	1.72	2.2	2.74
Mistral-8x7B-instruct-v0.1	4.67	6.28	7.18	6.52	4.75	2.97	3.22	3.33	3.17	4.68
DeepSeek-67B	5.17	6.09	7.23	6.42	3.85	2.97	3.67	3.62	3.58	4.73

Table 16: Subjective evaluation results on the dev set of the comparison-based feedback critique dimension.

Model	Test						Dev					
	Translation	Chat	QA	Summary	Harm.	Avg.	Translation	Chat	QA	Summary	Harm.	Avg.
<i>Closed-source LLMs</i>												
GPT-4	7.8	7.82	7.65	7.78	7.4	7.69	7.71	8.08	6.82	7.9	7.2	7.54
Gemini-Pro	7.1	7.59	7.35	7.39	8	7.49	-	-	-	-	-	-
Claude	7.52	7.25	7.58	8.02	8.22	7.72	7.32	7.62	7.45	7.7	7.75	7.57
GPT-3.5-turbo	7.38	7.62	7.58	7.32	7.84	7.55	7.18	7.28	6.95	7.5	7.03	7.19
PaLM	4.58	5.88	6.68	5.87	7.43	6.09	6	6.1	6.72	5.48	7	6.26
GLM4-no-tools	7.8	8.2	7.98	8	8.5	8.1	-	-	-	-	-	-
ErnieBot Pro	7.95	7.4	6.92	7.75	7.98	7.6	-	-	-	-	-	-
Baichuan2 Turbo	7.69	7.3	7.18	7.62	8.45	7.65	-	-	-	-	-	-
Qwen-Max	8.05	8.25	7.95	8.07	8.75	8.21	-	-	-	-	-	-
MiniMax-abab5	6.92	6.68	6.22	6.65	7.58	6.81	-	-	-	-	-	-
<i>Open-source LLMs (6B-7B)</i>												
InternLM2-7B-Chat	6.25	6.78	7.21	7.05	8.55	7.17	6.18	6.86	7.23	6.8	7.7	6.95
ChatGLM3-6B	3.05	4	5.72	4.92	7.78	5.09	3.75	4.42	5.17	5.42	6.47	5.05
Yi-6B	4.1	2.98	4.47	4.88	5.3	4.35	4.4	4.07	5.25	4.87	4.88	4.69
DeepSeek-7B	5.32	5.32	6.48	6.12	7.05	6.06	5.32	5.62	6.55	6.2	7	6.14
Baichuan2-7B-Chat	5.8	5.38	5.4	5.42	5.38	5.48	5.84	5.03	5.97	5.95	6.08	5.77
Qwen-7B-Chat	5.3	5.28	6.88	6.48	7.98	6.38	5.82	5.68	6.85	5.85	7.25	6.29
InternLM-7B-Chat	1.85	2.78	4.58	4.18	5.98	3.87	2.31	2.98	4.28	3.88	5.05	3.7
Llama2-7B-Chat	3.78	6.5	6.55	5.86	8.6	6.26	4.9	3.55	5.37	3.98	4.75	4.51
Vicuna-7B-v1.3	3.42	5.42	5.58	5.82	7.82	5.61	3.63	5.25	6.27	5.27	6.72	5.43
Mistral-7B-instruct-v0.2	5.45	7.02	7.35	7.7	8.48	7.2	5.47	7.07	7.43	7.33	7.88	7.04
<i>Open-source LLMs (13-20B)</i>												
InternLM2-20B-Chat	6.41	7.5	7.6	7.28	8.6	7.48	6.62	7.22	7.78	6.68	7.97	7.25
Qwen-14B-Chat	7.22	6.45	7.08	7.22	8.3	7.25	6.91	6.53	7.55	6.88	7.4	7.05
Baichuan2-13B-Chat	6.35	6.68	6.78	6.75	7.42	6.8	6.67	5.75	6.7	6.45	7.2	6.55
InternLM-20B-Chat	3.75	4.55	5.03	4.97	7.72	5.2	3.53	4.52	5.08	5.87	5.87	4.97
Llama2-13B-Chat	5.45	7	7.18	7.18	8.75	7.11	4.79	6.6	7.32	5.1	7.76	6.31
Yi-34B	6.12	5.08	5.82	6.32	7.9	6.25	6.57	5.3	5.75	6.52	6.92	6.21
Vicuna-33B-v1.2	5.1	7.25	6.8	7.58	7.9	6.93	4.68	6.8	7.22	7.02	7.35	6.61
WizardLM-13B-v1.2	5.31	6.22	6.4	5.89	8.35	6.43	5.39	6.32	6.56	6.14	7.22	6.33
<i>Open-source LLMs (> 70B)</i>												
Qwen-72B-Chat	7.16	6.88	7.25	7.62	8.35	7.45	6.95	7.22	7.22	7.64	7.45	7.3
Llama2-70B-Chat	5.58	7.2	6.85	7.18	8.72	7.11	3.33	5.84	6.72	3.96	7.85	5.54
Mixtral-8x7B-instruct-v0.1	5.18	7.88	7.5	7.5	8.6	7.33	5.07	7.42	7.34	7.68	7.95	7.09
DeepSeek-67B	7.02	7.2	6.72	7.2	8.36	7.3	6.72	7.17	6.6	6.92	7.1	6.9

Table 17: Subjective evaluation results on the test and dev set of the correction critique dimension. Due to the cost limitation, we do not provide the experimental results on these closed-source API-based LLMs: GLM4-no-tool, ErnieBot-Pro, Baichuan2 Turbo, Qwen-Max, MiniMax-abab5.

Model	Test					Dev				
	MathCoT	MathPoT	CodeExec	CodeNE	Avg.	MathCoT	MathPoT	CodeExec	CodeNE	Avg.
<i>Closed-source LLMs</i>										
GPT-4	50	62	83.33	83.33	69.67	40	74	80.95	75.61	67.64
Gemini-Pro	34	46	50	46.67	44.17	-	-	-	-	-
Claude	50	30	66.67	53.33	50	36	42	68.25	53.66	49.98
GPT-3.5-turbo	42	54	83.33	76.67	64.00	34	72	71.43	70.73	62.04
PaLM	25	31.25	16.67	40	28.23	30	42	33.33	24.39	32.43
GLM4-no-tools	40	56	73.33	73.33	60.67	-	-	-	-	-
ErnieBot Pro	43.75	56.25	83.33	60	60.83	-	-	-	-	-
Baichuan2 Turbo	38	48	56.67	46.67	47.34	-	-	-	-	-
Qwen-Max	56.25	50	60	55.56	55.45	-	-	-	-	-
MiniMax-abab5	32	46	53.33	43.33	43.67	-	-	-	-	-
<i>Open-source LLMs (6B-7B)</i>										
InternLM2-7B-Chat	28	20	50	46.67	36.17	24	32	55.56	43.9	38.87
ChatGLM3-6B	36	16	40	30	30.50	30	26	36.51	21.95	28.62
Yi-6B	18	0	10	10	9.5	4	6	14.29	12.2	9.12
DeepSeek-7B	38	42	33.33	26.67	35	40	48	31.75	26.83	36.65
Baichuan2-7B-Chat	25	12.5	16.67	0	13.54	30	34	20.63	19.51	26.04
Qwen-7B-Chat	38.00	28.00	30.00	33.33	32.33	34	40	25.4	29.27	32.17
InternLM-7B-Chat	4	4	10	10	7	0	12	9.52	7.32	7.21
Llama2-7B-Chat	0	0	0	0	0	0	2	1.59	0	0.90
Vicuna-7B-v1.3	20	24	30	16.67	22.67	20	28	15.87	14.63	19.63
Mistral-7B-instruct-v0.2	38	38	40	36.67	38.17	16.67	25	37.5	27.27	26.61
<i>Open-source LLMs (13-20B)</i>										
InternLM2-20B-Chat	44	38	60	60	50.5	30	52	61.9	56.1	50
Qwen-14B-Chat	40	32	46.67	33.33	38	34	48	42.86	26.83	37.92
Baichuan2-13B-Chat	26	36	33.33	30	31.33	16	42	28.57	24.39	27.74
InternLM-20B-Chat	40	18	10	6.67	18.67	40	18	20.63	7.32	21.49
Llama2-13B-Chat	30	4	0	0	8.50	24	2	0	0	6.5
Yi-34B	8	6	23.33	6.67	11	6	6	19.05	17.07	12.03
Vicuna-33B-v1.2	32	40	20	30	30.5	11.11	30	25	13.64	19.94
WizardLM-13B-v1.2	32	36	23.33	6.67	24.50	32	50	25.4	7.32	28.68
<i>Open-source LLMs (> 70B)</i>										
Qwen-72B-Chat	46	46	66.67	60	54.67	36	60	65.08	51.22	53.08
Llama2-70B-Chat	20	0	0	6.67	6.67	20	6	0	7.32	8.33
WizardLM-70B-V1.0	0	6	10	10	6.5	-	-	-	-	-
Mixtral-8x7B-instruct-v0.1	42	38	46.67	46.67	43.34	26	46	49.21	51.22	43.11
DeepSeek-67B	36	54	66.67	63.33	55	32	72	63.49	56.1	55.90

Table 18: Objective evaluation results on the test and dev set of the correction dimension. Due to the cost limitation, we do not provide the experimental results of following closed-source API-based LLMs on dev set: GLM4-no-tool, ErnieBot-Pro, Baichuan2 Turbo, Qwen-Max, MiniMax-abab5.